

A Synchronisation Facility for a Stream Processing Coordination Language


Anna Tikhonova

28 January 2015

AstraKahn

A program is a network of

- Boxes – **not** analysable
 - Synchronisers
- connected by
- Channels
 - Wiring patterns
 - Static
 - Dynamic (serial replication)



Objectives

A dedicated language of synchronisers:

- Implement a compiler
 - Transition analysis
 - Data transformation analysis

Output from the serial replication pipeline:

- Explore the role of synchronisers in the serial replication

Synchroniser

- State machine
 - Several input and output channels
 - Non-deterministic
 - Dedicated language (analysable)
- Store variables
 - Save messages
- State variables
 - Fold regular transition matrices

The Language (examples)


```
synch zip2 (a, b | c) {  
  store ma, mb;  
  start {  
    on:  
      a {  
        ma = this;  
        goto s1;  
      }  
      b {  
        mb = this;  
        goto s2;  
      }  
  }  
  s1 {  
    on:  
      b {  
        send ma || this => c;  
      }  
  }  
  s2 {  
    on:  
      a {  
        send this || mb => c;  
      }  
  }  
}
```

```
synch counter (a | b) {  
  state int(2) c = 0;  
  start {  
    on:  
      a & [c < 3] {  
        set c = [c + 1];  
      }  
      a & [c == 3] {  
        set c = [0];  
        send this => b;  
      }  
  }  
}
```

Term system

- Message Definition Language (MDL)
 - Basic terms
 - Constructors
- Examples:
 - Variable $\$t$
 - Record $\{ \ x:\$x, \ y:\$y, \ z:\$z \ }$
 - Choice $(: \ v:\$v, \ w:\$v :)$

Passports



`a.(x, y || t)`
`send this => c`

`a ~ (: uniq:{x:$x, y:$y | $t} :)`
`c is the same as a`

`b.?w(r, k)`
`send ?z(x:[r]) => d`

`d ~ (: z:{x:Int} :)`

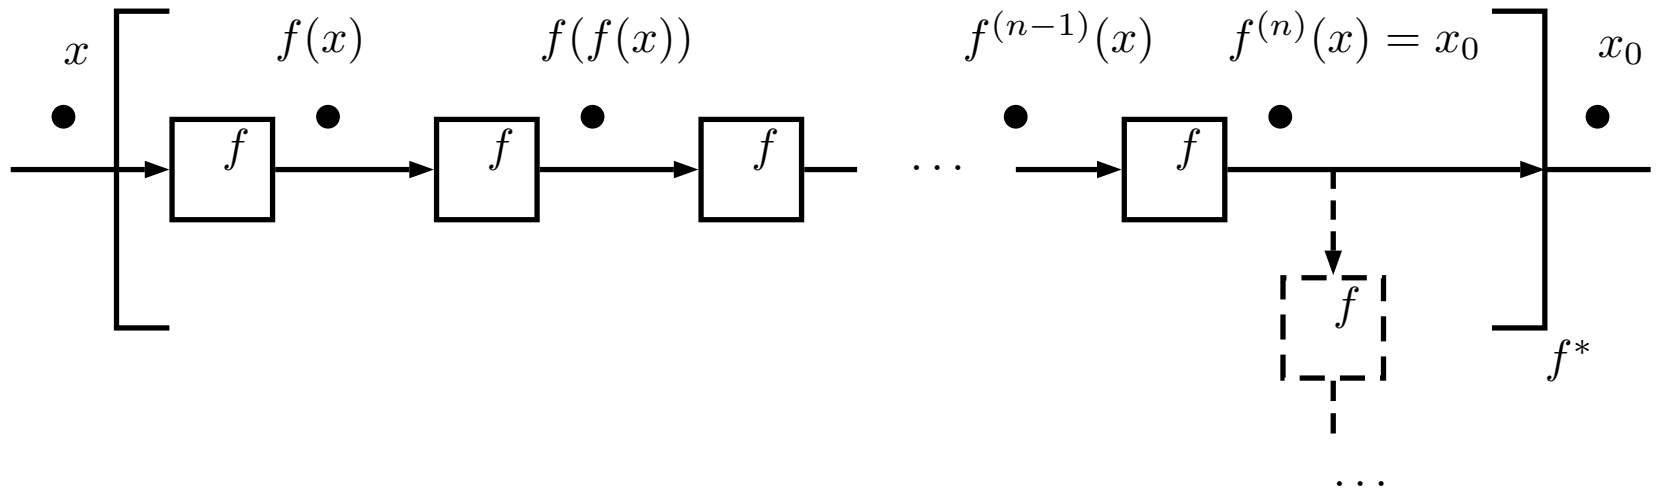
`b.?v`

`b ~ (: v:$v, w:{r:Int, k:$k} :)`

The compiler

- Lexical and syntax analysis
- Static analysis
 - Semantic checking
 - Type checking
- Runtime data structure
- Passport

Serial Replication




x_0 is the fixed point message

An inactive replica transmits the fixed point message
without change

Forward Fixed Point


Every synchroniser on the fixed-point path transmits the fixed-point message without change



The fixed-point condition of \mathbf{N} is

$$p_1 \wedge (p_2^1 \vee p_2^2) \wedge p_3$$

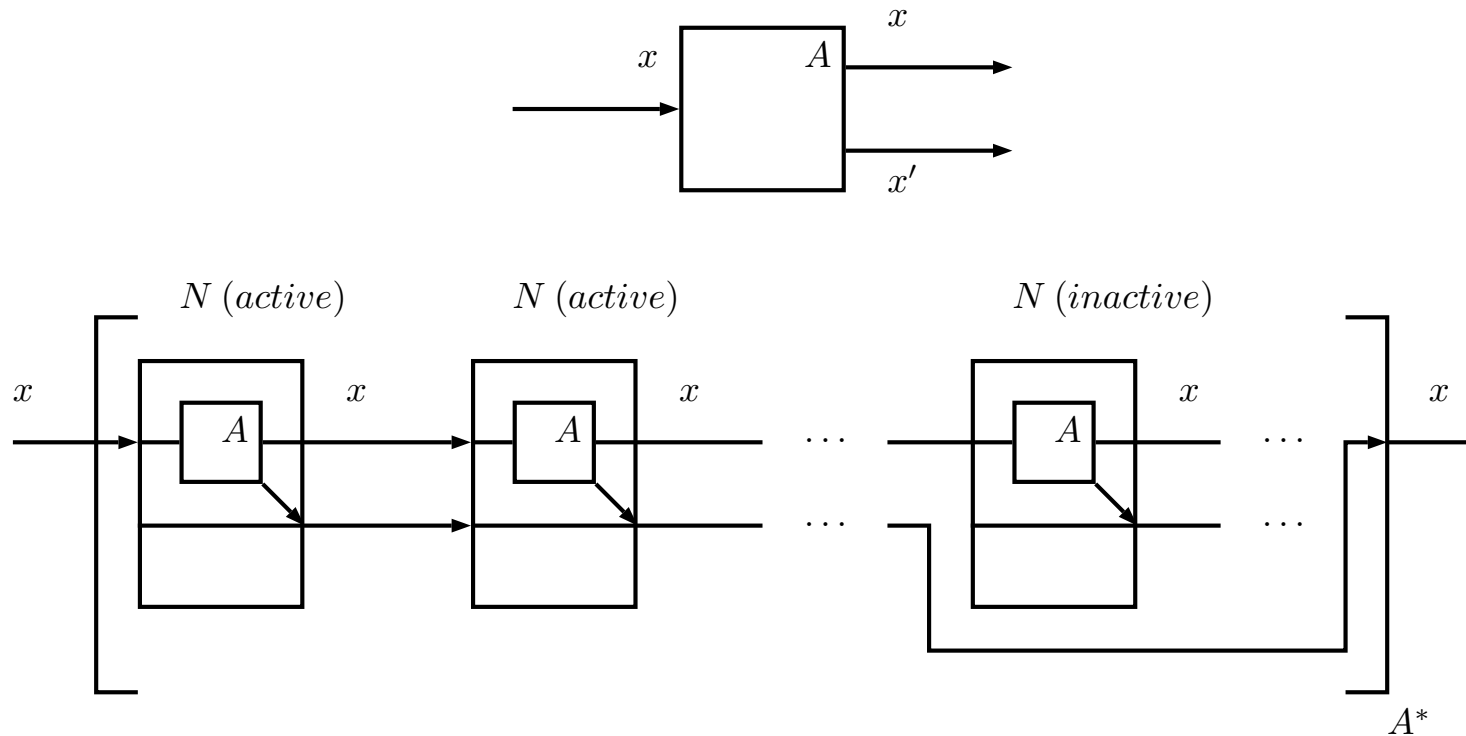
Output from the Serial Replication



Synchroniser $\mathbf{S}^{(j)}$, $j=0,m$ checks if the fixed point condition extracted from the \mathbf{j} -th synchroniser on the fixed point path is satisfied:

- **Yes** – the message is sent to $\mathbf{S}^{(j+1)}$
- **No** – the message is sent to the next replica of \mathbf{N}

Another Approach



Channel x' is auxiliary

Reverse Fixed Point

- prevents an unnecessary cascade
- is a state of a network in which it transmits any message without change
- A synchroniser can have a reverse fixed point
- A replica has a reverse fixed point if:
 - The replica has the fixed-point path
 - Every synchroniser on the fixed-point path has a reverse fixed point (possibly a closed subset of states)

Fixed Point Detection

- The fixed-point path search
 - DFS-based
 - Shared for both types of fixed point
- Fixed point detection in a synchroniser
 - Forward fixed point
 - Reverse fixed point

Conclusion

- The compiler
 - Transition analysis
 - Data transformation analysis
- Analysis of the fixed point approach
 - Another approach suggested
- Future directions
 - Flow inheritance
 - Statistical properties