

A Synchronisation Facility for a Stream Processing Coordination Language

Anna Tikhonova

A thesis submitted to the University of Hertfordshire
in partial fulfilment of the requirements of the degree of

Master of Science by Research

January 2015

Abstract

Contents

Abstract	i
1 Introduction	1
1.1 Motivation and Contribution	1
1.2 Outline	2
2 Foundations and Related Work	4
2.1 Coordination Programming	4
2.2 Stream Processing	6
2.3 Coordination in Streaming Networks	7
2.4 Summary	14
3 <i>AstraKahn</i> Synchroniser	15
3.1 Mathematical Model	15
3.2 The Language of <i>AstraKahn</i> Synchronisers	19
3.3 Execution Order of Synchroniser	25
3.4 The Implementation of the <i>aksync</i> Compiler	27
3.5 Discussion and Future Work	38
4 Serial Replication in <i>AstraKahn</i>	40
4.1 <i>AstraKahn</i> Approach to Serial Replication	40
4.2 Forward Fixed Point	42
4.3 Reverse Fixed Point	49
5 Conclusion	53
5.1 Summary	53
5.2 Future Work	54
A The Syntax of the <i>AstraKahn</i> Synchroniser	55
A.1 Full Grammar	55
A.2 Keywords, Reserved Words and Punctuation	55
A.3 The Abstract Syntax Tree of the Synchroniser	55
Bibliography	59

Bibliography

59

Chapter 1

Introduction

For years, processor manufacturers have delivered increases in clock rates. While manufacturing technology still improves, physical limitations of semiconductor-based electronics have become a major concern of design. In order for the processors to continue to improve in performance, multi-core design has become necessary. However, increasing the number of processor cores does not provide automatic benefits for legacy applications. Parallelisation of the performance-demanding parts of an application is difficult because computation concurrency management are mixed within the application code. One approach to address this issue is to present the application as a collection of independent components and specify their communication in a so-called coordination language, thus separating the concerns of computation, coordination and concurrency management.

The concepts of the new coordination language *AstraKahn* are described in [1]. The language defines the coordination behaviour of asynchronous stateless components (*boxes*) and their orderly interconnection via stream-carrying channels with finite capacity. *AstraKahn* structures the interconnection using a fixed set of wiring primitives, viz. serial and parallel composition, wrap-around connection and serial replication. Boxes are connected to the network with one or two input channels and one or more output channels. A stateless box does not synchronise data on its input and output channels; to this end, *AstraKahn* provides a synchronisation facility called *synchroniser*. In order to deal with the issue of application progress, *AstraKahn* attempts to provide an automatic resource and concurrency management based on communication demand.

1.1 Motivation and Contribution

At the moment, the *AstraKahn* project is on the initial stage and there exists no software implementation at all. In order to carry out research towards automatic resource and

concurrency management, an execution environment for *AstraKahn* applications must be developed. In brief, such an environment includes a compiler to generate an intermediate representation of the application source code and a runtime system to interpret the representation under the input data. Before the execution environment can be implemented, the main concepts of *AstraKahn* must be well-established.

This thesis focuses on the implementation of synchronisers and their role in the serial replication wiring pattern. Synchronisers are programmed in a dedicated language that is described in [1]. We provide some minor syntax improvements and explain how each language construct should be used in a detailed guide. An *AstraKahn* synchroniser has non-deterministic behaviour. We give an execution algorithm which defines the ordering of non-deterministic choices made by the synchroniser, and is a basis for the synchroniser runtime. We implement the language compiler that generates the data structure to be interpreted by the runtime, and the communication passport of the synchroniser. The compiler performs some semantic and type checking and reports errors in the source code as well.

In *AstraKahn* the output from the serial replication pipeline is defined using the concept of fixed point. A fixed point is a message that is not changed after it has been processed by a newly created replica. In order to detect fixed point messages, *AstraKahn* must be provided with a pattern that matches all of them. We show how this pattern can be embedded into the operand network of the serial replication combinator, so that the programmer does not have to specify it explicitly within the *AstraKahn* application code. Since the original approach to the output from the serial replication is quite complicated, we provide a simple alternative. In order to suppress the growth of the replica chain, *AstraKahn* introduces a reverse fixed point, which is a state of a replica that is known not to change an input connection anymore and thus can be optimised out. We give formal definitions of both kinds of fixed point and provide the algorithms for the *AstraKahn* compiler to detect them.

1.2 Outline

The remainder of the thesis is as follows. Chapter 2 introduces *AstraKahn* and presents some theoretical prerequisites in coordination programming and stream processing. *AstraKahn* is compared with a recent component system example from each field.

In Chapter 3 the implementation of the synchroniser is described in detail. We describe a synchroniser mathematically in order to explain some facilities it needs. Also, the

chapter includes the language guide, the execution algorithm and the implementation of the compiler.

In Chapter 4 the machinery behind the serial replication in *AstraKahn* and the role of synchronisers in it is explained.

Chapter 5 concludes the thesis, providing directions for further research.

Chapter 2

Foundations and Related Work

In this chapter we provide relevant theoretical background in coordination programming and stream processing. We pick a recent component-system example from each field. Then, we describe the combined approaches to coordination programming and stream processing, implemented in S-NET and *AstraKahn* and explain the concepts behind *AstraKahn* in more detail. We sum up the chapter with a comparison of the four components systems, including their approaches to synchronisation.

2.1 Coordination Programming

The coordination paradigm offers a promising way to address some issues related to the development of efficient parallel systems. Programming a parallel system can be seen as a combination of two activities: the actual computing part comprising a number of processes that manipulate data and a coordination part that is responsible for communication between the processes.

In the main, coordination is managing dependencies between components. Since the computation is completely separated from the coordination, the processes that comprise the former are seen as black boxes. The programming languages used to write the computational code do not play an important role in setting up the coordination scheme.

Existing coordination models¹ are described in details in the survey [2] by G. Papadopoulos and F. Arbab. They argue that these models fall into two major categories of coordination programming, namely either data-driven or control-driven.

¹A coordination model encompasses entities being coordinated, a means of coordination and a semantic framework

The main characteristic of data-driven coordination models is that the coordination code is mixed with the process definition. A data-driven coordination language typically offers several coordination primitives which are intertwined with the purely computational code. Many data-driven coordination models have evolved around the notion of shared dataspace. The shared dataspace plays a dual role, being a global data repository and an interprocess communication system. The processes communicate by writing to the shared dataspace and retrieving data from it. Historically the first member of this family is LINDA [3]. Strictly speaking, not all data-driven coordination models follow the above pattern of coordination. Some of them use a message-passing based mechanism (MPI, [4]).

Opposite to the data-driven coordination model, the control-driven coordination achieves almost complete separation of concerns between computation and coordination. This is usually achieved by defining a special language that offers facilities for controlling synchronisation, communication, creation and termination of computing components. One of a contemporary members of this family is REO [5]. In REO computational components communicate via complex coordinators, or *connectors*. An undirected channel is an atomic connector in REO. Channels are typed, however, no fixed set of types is assumed. The channel type defines the behaviour of the channel with respect to data. A list of common types is as follows:

- **Sync** A channel of type *Sync* atomically gets data from the input and propagates it to the output
- **Lossy Sync** Same as *Sync*, but data can be lost if the output is not ready to accept data
- **FIFO(n)** A channel of type *FIFO(n)* gets data from the input, temporarily stores it in an internal buffer of size n , and propagates it to the output whenever it is ready to accept the data
- **Sync Drain** A channel of type *Sync Drain* atomically gets data from both inputs and loses it
- **Filter(c)** A channel of type *Filter(c)* atomically gets data from the input and propagates it to the output if the filter condition c is satisfied. Otherwise, the data are lost.

Channels are connected with *nodes*. Nodes have fixed merger-replicator behaviour: the data of one of the incoming channels is propagated to all outgoing channels, without storing or altering it. If multiple incoming channels can provide data, the node makes a

nondeterministic choice among them. A complex connector in REO is represented as an undirected graph of channels and nodes. C. Baier et al. propose *constraint automata* as an operational model for component connectors in REO [6].

2.2 Stream Processing

A stream processing system is a system comprised of a collection of isolated processes that compute in parallel and communicate data solely via static channels. The processes are usually divided into three classes: sources that create data for the system, filters that perform some computation, and sinks that pass data from the system. Stream processing systems are usually visualised as directed graphs.

An overview of stream processing is given in the survey by R. Stephens [7]. Stephens identifies that the earliest type of stream processing systems is dataflow. In the first dataflow programming language LUCID [8], each variable is represented as an infinite stream of values. Computation is carried out by defining transformation functions that process such streams. Lucid is possibly the first language to introduce the idea of filter.

A significant result for concurrency engineering is Kahn's work [9], which outlines the semantics of a simple parallel programming language. Kahn suggests a distributed model of computation where a group of deterministic sequential processes communicate via unbounded FIFO channels under the following assumptions:

- Channels are the only way for processes to communicate
- Channels transmit messages within a finite time
- At any given time a process is either performing computation or waiting for messages on a specific input channel.

Kahn proved that the output of the resulting process network is deterministic, i.e. it does not depend on the ordering of computations at different nodes. The model is commonly referred to as Kahn Process Network (KPN).

A Kahn process may have multiple input and multiple output channels. Reading from a KPN channel is blocking, i.e. a process that reads from an empty channel stalls and can only continue when the channel contains sufficient data. On the contrary, writing to a channel is non-blocking, and it always succeeds since the capacity of a KPN channel is unlimited. Processes cannot test an input channel for data availability without committing to consume the data. KPNs allow arbitrary wiring, i.e. a network may have feedback communication.

In KPNs the number of data elements a process might read from a channel or write to a channel is not restricted. In synchronous dataflow (SDF, [10]) the consumption and production rates of a process are fixed. A recent SDF language is STREAMIT [11]. The basic unit of computation in StreamIt is a user-defined single-input single-output (SISO) block called a filter. The filter can communicate with neighbouring blocks via FIFO channels. STREAMIT structures an application using the following primitives:

- *Pipeline* specifies sequential composition of filters
- *SplitJoin* specifies parallel composition of filters
- and *FeedbackLoop* provides a way to create loop constructs in a streaming network.

A STREAMIT program is a hierarchical composition of these constructs.

Thanks to the single-input and single-output restriction, a filter does not need to synchronise data on multiple input channels and to split result between output channels.

2.3 Coordination in Streaming Networks

2.3.1 S-Net

S-NET [12] is a declarative coordination language based on stream processing. It defines the behaviour of stateless asynchronous components (boxes) that interact with each other in a streaming network. Boxes are written in conventional languages that are subject to contract with S-Net. They execute fully asynchronously, i.e. a box may consume data as soon as it is available from the input stream. Moreover, boxes are SISO, therefore S-NET achieves a near-complete separation of concerns between communication and computation.

Streaming networks are expressed in a hierarchical manner using a fixed set of five combinators, viz. serial composition, parallel composition, serial replication, parallel replication and feedback loop. Three of the five combinators have non-deterministic versions that permit arbitrary reordering of input streams.

Data on streams are organised as variant records of label-value pairs. S-NET provides a special facility, called *synchrocell*, that merges one or more records into a single one. A synchrocell maintains an internal state in order to keep incoming records which match one of the patterns until all patterns have been matched. Then the records are merged into a single one and sent to the output stream. A synchrocell provides storage for

one record of each pattern, and records with an already matched pattern are forwarded directly to the output stream. After sending the result of the merging on, the synchronocell serves as an identity function, forwarding all incoming records to the output. In order for the synchronocell to merge continuously, a serial replication network combinator must be applied to it.

2.3.2 *AstraKahn*

AstraKahn is an attempt to provide a component system with automatic concurrency management. *AstraKahn* defines the coordination behaviour of fully asynchronous components (boxes) and their orderly interconnection via stream-carrying bounded FIFO channels. In *AstraKahn* data on streams are organised as sequences of messages. Each message conforms to one or more statically known formats.

AstraKahn provides a facility for stream synchronisation in the form of a special component called a *synchroniser*. The behaviour of the synchroniser is not fixed; instead, it is defined in a dedicated language that is a part of *AstraKahn* paradigm. An *AstraKahn* box generally is not SISO. Typically it has a single input channel, however, the number of output channels, although statically known, is not restricted. Similar to S-NET, boxes are stateless, hence they do not synchronise data; this work is done by synchronisers.

AstraKahn structures streaming networks using a total of four combinators, namely: the serial connection, the parallel connection, the wrap-around connection and the serial replication. Network combinators may take either boxes or networks as their operands, hence the network construction is an inductive process.

Note that *AstraKahn* does not have the parallel replication combinator which exists in S-NET. Thanks to the statelessness of boxes, the parallel replication can be implemented with the serial replication combinator. The routing of messages is managed by an array of synchronisers that are indexed within the declared limits.

In the following sections the concepts of *AstraKahn* are explained in more detail.

2.3.2.1 Channels

Channels in *AstraKahn* are named FIFO queues with a limited capacity. A channel carries a segmented stream that consists of message sequences and those may in turn consist of sequences in their own right. In order to mark the beginning and end of a sequence, *AstraKahn* supports a special kind of message called a segmentation mark.

Segmentation marks can be thought of as brackets. *AstraKahn* requires that a stream of message sequences that flows through a channel has a static bracketing depth. Therefore, each message on a given channel is found between the same number of brackets. The sequence of messages starts with a certain number of opening brackets and ends with the same number of closing brackets. Within the sequence brackets can occur only in the following combination:

$$\underbrace{)\dots)}_k \underbrace{(\dots(}_k,$$

where $k \leq d$, and d is the number of opening brackets in the beginning of the stream. This combination constitutes the segmentation mark σ_k . The bracketing depth $d \geq 0$ is a static characteristic of a channel².

2.3.2.2 Boxes

Boxes are the atomic building blocks of *AstraKahn* networks that perform the computation. An *AstraKahn* box is deterministic in the sense that for every partial input stream it produces a deterministic output stream³.

Conceptually, boxes can be specified in any conventional programming language; however, they are subject to a contract that defines acceptable behaviour for boxes. Any guarantees that *AstraKahn* offers are subject to the fulfilment of the contract on behalf of all the boxes. The interface between a box and the *AstraKahn* runtime system is defined by the *AstraKahn* Box-API for each supported box language.

AstraKahn declares seven box categories with respect to their algebraic properties and effect of channel segmentation⁴:

Transductor A transductor has one input channel and one or more output channels and responds with no more than one output message on each of its output channels. Segmentation marks are passed on to all the output channels of the box, bypassing the box code.

Inductor An inductor has one input channel and one or more output channels and responds to a single message from the input channel with a sequence of messages

²Indeed, the bracketing depth of a channel that would carry the stream of message lists

$$(((\underbrace{a)}_{\sigma_1})(\underbrace{b)}_{\sigma_2}))((\underbrace{c)}_{\sigma_1})(d)))$$

is 3

³For a function $f(x) : \mathcal{I} \rightarrow \mathcal{O}$, where \mathcal{I} is the totality of $f(x)$ input streams and \mathcal{O} is the totality of $f(x)$ output streams, $\forall p \in \mathcal{I} \wedge \forall t : p \cup t \in \mathcal{I} : f(p \parallel t) = f(p) \parallel f(t)$

⁴The box code does not see the segmentation marks; *AstraKahn* deals with them all by itself

on each of its output channels. Before the input stream is passed to the inductor, each σ_k in it with $k > 0$ is replaced by σ_{k+1} , and a σ_1 is inserted between every two consecutive data messages. Segmentation marks are bypassed from the input to all the output channels by the coordinator when encountered at the input of the inductor.

Reducer A reducer implements the reduction operation for a list of input messages. The reducers can have more than one output channel with one of them reserved for the results of the reduction. *AstraKahn* classifies reducers by the number of input channels and properties of the reduction operation they implement. There exist five classes of reducer:

Dyadic ordered A dyadic ordered reducer has two input channels. The first input channel is reserved for the initial value. The reduction operator is applied to the messages in the order that they arrive on the second input channel

Dyadic unordered Same as dyadic ordered except that the reduction operator can be applied to the messages on the second channel in any order without affecting the result

Monadic ordered and monadic unordered Same as dyadic reducers except monadic reducers have one input channel. A monadic reducer is only started when two messages are received

Monadic segmented A monadic reducer recursively processes an input list of messages that can be segmented into arbitrary sublists until the list is reduced to a single message

2.3.2.3 Synchronisers

Synchronisers are non-deterministic finite state machines for joining messages and sending them on to the output channels. *AstraKahn* provides synchronisers with memory for storing received messages.

A synchroniser can have any number of input and output channels. Unlike boxes, synchronisers maintain an internal state and generally accept messages from an input channel in certain states, while in another state the channel may be blocked until a state transition brings the synchroniser to a state in which messages from the channel are accepted.

A synchroniser can also compute trivial extensions for messages. For example, it can append a labeled integer value to a message. It also detects segmentation marks in an

input stream and can change the segmentation of the stream by sending segmentation marks to the output channels.

The state transitions of a synchroniser can depend on the content of the current message but never on that of a stored one. In order for the synchroniser to read values from the current message, it is matched with a pattern specified within the triggering condition of the transition. In addition, the triggering condition may check the matched values if they are known to be integers. If the message was matched and the integer values satisfy the condition, then the transition fires.

The act of sending a message to an output channel is associated with a transition. Once the transition is known to fire, the synchroniser computes the message extensions, combines all the parts of the message together and sends it on to the output channel.

AstraKahn provides a dedicated language of synchronisers.

2.3.2.4 Network Composition

The construction of streaming networks in *AstraKahn* is hierarchical: components are wired into a subnetwork, which in turn can act as a component in a larger network, etc. In order to wire the components, *AstraKahn* provides a set of wiring patterns sufficient to achieve arbitrary wiring. Each pattern identifies input/output channels of the operand(s) with one another and with the input/output channels of the result.

Three patterns are static, applicable to one or two operands:

Serial connection applies to two operands. All outputs of the first operand are wired to identically named inputs of the second operand if they exist. The rest of the inputs and outputs contribute to the input/output sets of the resulting network.

Parallel connection applies to two operands. Two operand networks are placed side by side without connection and their input and output channels form the input and output channel sets of the resulting network.

Wrap-around connection applies to a single operand. Each output channel of the operand that matches an input channel by name is wired to it with a special wrap-around channel, thus completing a cyclic connection. In order to avoid deadlocks, *AstraKahn* does not limit the capacity of wrap-around channels; their capacity is only limited by the amount of memory available for the queues in the system⁵.

⁵if there is no memory available, the application crashes

These three patterns are sufficient to achieve arbitrary wiring of the components.

The fourth pattern – **serial replication** – replicates the single operand network infinitely and wires up the replicas with the serial connection. In implementation, actual replication is demand-driven: the chain of replicas is extended dynamically. Messages are extracted from the chain and sent to the output channel when they satisfy the fixed point condition, see Chapter 4.

2.3.2.5 The Type System of *AstraKahn*

The type system of *AstraKahn* is based on the Message Definition Language (MDL, [1]) which is a language of abstract terms that are built recursively from the ground up. Structurally the terms are symbolic trees with the following kinds of leaf:

Symbol, Number, String terms represent a certain finite quality

Variable term ranges over terms

Flag is a boolean variable that only occurs in a certain context.

Other terms are built recursively using the following types of constructor:

Tuple is a sequence of terms in linear order

List is an extensible sequence of terms in linear order

Record is an extensible collection of label-term pairs

Choice is an extensible collection of alternative labeled terms

Switch is a collection of guarded terms that represents exactly one of them depending on the value of the boolean guards

Data on streams are organised as sequences of messages defined by a **choice**. If a choice is known to carry a single record, this record is labeled *uniq*.

An *AstraKahn* component, either a box or a synchroniser, is both a consumer and a producer for some other components in the network. Hence to guarantee the static correctness of a connection, the subtyping relation between the consumer's input and the producer's output types must be satisfied. In order to check the static correctness over the network, a component can be abstracted with respect to its data-transformation behaviour as an implicative statement $p \Rightarrow P$, called a passport, where p is the type

of the input message and P is the type of the output message. During the check, the *AstraKahn* compiler extracts the topology of the network, forms the subtyping relations between the passports and performs constraint solving in order to instantiate all term variables. If the constraint system is satisfiable, then the whole program is consistent and type correct.

2.3.2.6 Structure of *AstraKahn*

AstraKahn is motivated by a KPN, which is a theoretical model and its properties are only available under an interpretation with unlimited resources. The intention of *AstraKahn* is to refine and structure KPNs in order to provide a component system with an automatic resource and concurrency management.

AstraKahn provides the following refinements to KPNs:

- *Explicit management of state in the framework of coordination program* gives rise to more usable parallelism. Computational components are stateless and the coordination logic is expressed in the network wiring and synchronisers
- *Automatic resource and concurrency management* provides a parallelisation mechanism based on communication demand
- *Separation into independent layers communicating by means of interfaces* addresses the standard engineering agenda of abstraction, encapsulation and hierarchical development.

As the result of the refinement, *AstraKahn* is presented as a paradigm with the following layers:

- A Topology and Progress Layer (TPL) that defines:
 - Classes of boxes, their algebraic properties and their effects on channel segmentation
 - The language of synchronisers
 - The wiring patterns and the subnetwork encapsulation facility
 - The automatic resource and concurrency management strategy
- A Constraint Aggregation Layer (CAL) that ensures type safety all over the network given the data constraints supplied by each component
- A Data Instrumentation Layer (DIL) that manages data migration and concurrent access to objects in memory.

2.4 Summary

Earlier on we reviewed recent component systems based on coordination programming (REO) and stream processing (STREAMIT, S-NET), and described the approach to component coordination being developed within the *AstraKahn* project. The stream processing based approaches STREAMIT, S-NET and *AstraKahn* impose structuring on networks with fixed sets of combinators, while the coordination language REO only supports unstructured component connection. In REO, the computational components are connected into a network with complex connectors that are constructed of channels typed with respect to their synchronisation properties. Just like the REO's approach to data synchronisation, S-NET achieves a near-complete separation of concerns between computation and communication. However, in S-NET, a computational component's interface is restricted to SISO. Additionally, S-NET provides a stream synchronisation facility called synchrocell. *AstraKahn* does not introduce the restriction for computational components. Instead, *AstraKahn* provides a more complex stream synchronisation facility. Unlike the REO connector or the S-NET synchrocell, the *AstraKahn* synchroniser is able to process messages, e.g. read and change their content to some extent. With synchronisers, *AstraKahn* achieves a separation of concerns between computation and communication. Similar to S-NET, the computational components in STREAMIT are SISO. Moreover, STREAMIT is based on the synchronous dataflow model, where neighbouring components communicate synchronously.

In order to support dynamic reconfiguration of streaming networks, S-NET and *AstraKahn* require computational components to have no state. A heuristic scheduler that utilises positive and negative demands of the stream communication was developed for S-NET in [13]. The ability to dynamically reconfigure a streaming network opens opportunities for parallelisation. The long-term goal of the *AstraKahn* project is to provide a self-regulating concurrency mechanism based on communicational demand. STREAMIT does not require the components to be stateless; it relies on the static scheduling of the synchronous data flow programs. Reo is clueless about the components it coordinates; it focuses on the components connection, rather than on the components themselves. In other words, the problem of automatic parallelisation is not set for both STREAMIT and REO.

Chapter 3

AstraKahn Synchroniser

In this chapter we describe the source language for an *AstraKahn* synchroniser and its implementation. Prior to describing the language in details, we present a mathematical model of a synchroniser from [1] and outline sources of non-determinism in synchronisers.

The *aksync* compiler is integrated into the current *AstraKahn* runtime system prototype. It takes the source code of a synchroniser program, performs the syntactic and semantic analysis of it, builds a synchroniser object and generates the CAL passport of the synchroniser.

3.1 Mathematical Model

From the mathematical point of view a synchroniser is a pair

$$(\Phi, \Pi),$$

where $\Phi = (A, S, T)$ is a nondeterministic state machine with the alphabet of events $A \subseteq C \times P$, where C denotes the set of input channels and P the set of the predicates on channel messages¹. The set of abstract states in Φ is denoted as $S \supseteq \{s_0\}$ with the start state s_0 , $T : A \times S \rightarrow S$ is the transition matrix of Φ . The path functional $\Pi : S \times \Omega \rightarrow V^{(*)}$ defines the output of the synchroniser, where Ω is the set of output channels and V is the set of message values².

¹An event $(c, p) \in A$ represents the reception of a message on channel c that satisfies the predicate p

² $V^{(*)}$ denotes a set of all sequences from V

In state s_k the functional is based on the retrospective sequence of transitions from the most recent visit to the start state s_0 to s_k :

$$(s_0, c_0), (s_1, c_1), \dots (s_k, c_k),$$

where $c_i \in C$, $0 \leq i \leq k$ is the channel that caused the transition from the state s_i .

Let μ_i be the message received in the transition from the state s_i . Then

$$\Pi(s_k, \omega_m) = \psi_{\sqcap} \{ \mu_i \mid \rho_{ki}^m(s_i), 0 \leq i \leq k \}, \quad (3.1)$$

where ρ_{ki}^m is the selection predicate that defines Π , and ψ_{\sqcap} is the operator that coerces the messages in the operand set to their joint greatest subtype.

From the above, the synchroniser is fully defined by two functions:

1. The transition matrix T

The state machine can have a regular structure whereby many transitions can be defined at once by a formula with some limited-range integer variables. For example, a machine with 8 states could have a transition matrix defined thus: $S_{k \bmod 8} \rightarrow S_{k+1 \bmod 8}$. In order to be able to employ regular transition graphs, *AstraKahn* allows synchronisers to declare *state* variables.

Example: the counter synchroniser Counter sends every n -th message from its input channel to the output channel, other messages are disregarded. The transition diagram for the counter synchroniser for $n = 3$ is given in Figure 3.1.a.

The counter synchroniser is a pair (Φ, Π) , where

$$\begin{aligned} \Phi &= (A, S, T), \\ C &= (a), P = (true), A = C \times P = ((a, true)), \\ S &= (s_0, s_1, s_2), s_0 - \text{start state}, \\ T: &\frac{A \setminus S}{(a, true)} \begin{array}{c|c|c} s_0 & s_1 & s_2 \\ \hline s_1 & s_2 & s_0 \end{array} \\ \Pi : S \times \Omega &\rightarrow V^{(*)}, \\ \Omega &= (c), \\ V &= (a) \end{aligned}$$

An output message is emitted when a transition happens from the state s_2 . This state is reached in a single path:

$$\begin{aligned} W_0 &= ((s_0, a), (s_1, a), (s_2, a)) \Pi(s_2, c) = \psi_{\sqcap} \{ \mu_0 = a \mid \rho_{20}^c(s_0) = 0, \mu_1 = \\ &a \mid \rho_{21}^c(s_1) = 0, \mu_2 = a \mid \rho_{22}^c(s_2) = 1 \}, k = 1, i = 0, 1, 2 \end{aligned}$$

The state machine behind the counter has a regular structure, and for this synchroniser all its transitions can be defined with a single formula: $S_k \bmod 3 \rightarrow S_{k+1 \bmod 3}$.

Considering this, the transition matrix T would be:
$$\frac{A \setminus S}{(a, true)} \mid \frac{S_k \bmod 3}{S_{k+1 \bmod 3}}$$

Some possible transition diagrams of the counter synchroniser are given in Figure 3.1. The diagram 3.1.a represents the unrolled regular structure of the synchroniser. However, this representation is inconvenient when $n \gg 1$. The diagram can be folded using state variables. Two possible variants are shown in figures 3.1.b and 3.1.c. The state variable c acts as an induction variable in a while loop with the exit condition $c \geq 3$.

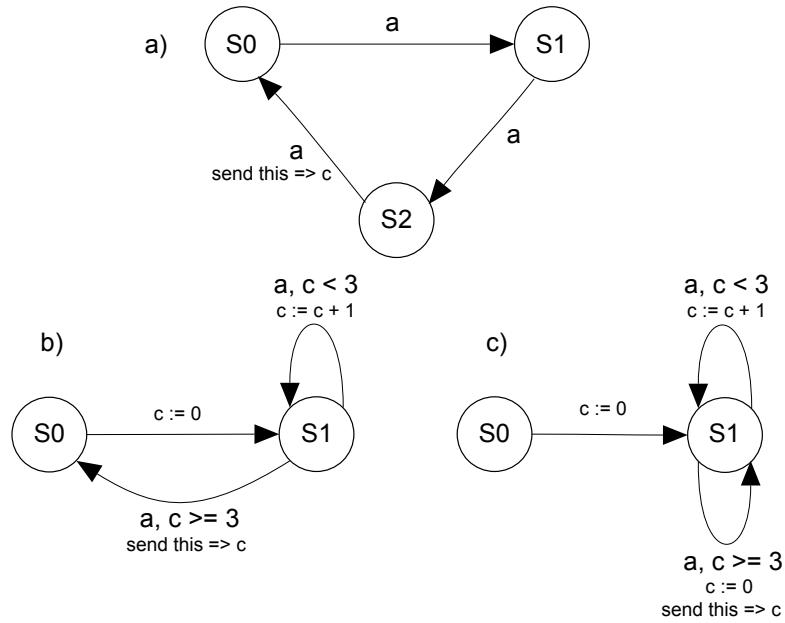


FIGURE 3.1: The transition diagrams of the counter synchroniser

2. The selection predicate ρ (see formula 3.1)

In a given state k for each output channel ω_m we note all i on which ρ_{ki}^m is true. Those message values must be saved in a previous state and recalled in state k . The functional Π can be implemented as function that retrieves all the saved messages at once; however, it is expected that the boolean vector $\omega_i = \rho_{ki}^m$ has only very few true elements. Consequently it is feasible to implement the storage mechanism that *AstraKahn* provides for synchronisers in the form of individual *store* variables. The type of a store variable is determined when a variable is assigned.

Example: the binary zip synchroniser Zip2 receives messages on its input channels and sends their concatenation to the output channel. In the resulting concatenation there is exactly one message from each input channel and those messages are combined for the output.

The zip2 transition diagram is given in Figure 3.2. The message received in the current transition is referred by a keyword *this*. *ma* and *mb* are the store variables associated with the input channels *a* and *b* respectively. The statement *send* indicates the sending of a message to an output channel.




FIGURE 3.2: The transition diagram of the zip2 synchroniser

Mathematically, the zip2 synchroniser is a pair (Φ, Π) , where

$$\Phi = (A, S, T),$$

$$C = (a, b), P = (true), A = C \times P = ((a, true), (b, true)),$$

$$S = (s_0, s_1, s_2), s_0 - \text{start state},$$

$A \setminus S$	s_0	s_1	s_2
$T: (a, true)$	s_1	s_1	s_0
$(b, true)$	s_2	s_0	s_2

$$\Pi : S \times \Omega \rightarrow V^{(*)},$$

$$\Omega = (c),$$

$$V = ((a, b))$$

An output message is emitted when a transition happens either from state s_1 or state s_2 . These states are reached in two paths:

$$W_0 = ((s_0, a), (s_1, b))$$

$$\Pi(s_1, c) = \psi_{\cap} \{ \mu_0 = a \mid \rho_{10}^c(s_0) = 1, \mu_1 = b \mid \rho_{11}^c(s_1) = 1 \}, k = 1, i = 0, 1$$

$$W_1 = ((s_0, b), (s_2, a))$$

$$\Pi(s_2, c) = \psi_{\cap} \{ \mu_0 = b \mid \rho_{20}^c(s_0) = 1, \mu_2 = a \mid \rho_{22}^c(s_2) = 1 \}, k = 2, i = 0, 2$$

So far we have given a formal definition of the synchroniser. Along with it we have introduced two important concepts for the synchroniser that are considered as a part of the synchroniser language. First is a message storage mechanism we call a store variable. Second is a mechanism for defining large regular transition matrices, which we call a state variable.

3.2 The Language of *AstraKahn* Synchronisers

The *AstraKahn* synchroniser is a finite state machine, therefore the basic building blocks of a synchroniser program are states and transitions. A state of a synchroniser is fully defined by the corresponding state of the finite state machine and the values of the state variables. A transition is the act of moving to another state which is initiated by a triggering event. A triggering event for the synchroniser transition is the arrival of a message to the associated channel. The message may be required to have a specific structure. In addition, a transition may be guarded by special conditions on state variable values. If the condition is satisfied the transition fires, otherwise it is cancelled.

Once a transition is known to fire, optional actions may be performed before the underlying state machine makes the move. These actions include changing the state and store variables and sending messages to output channels. In order to change the state and store variables, the synchroniser language provides state and store expressions over them.

This section gives an overview of the *AstraKahn* synchroniser programming language. The formal grammar of the *AstraKahn* synchroniser is provided in Appendix [A.1](#).

3.2.0.7 Program structure

A synchroniser program consists of a header followed by a body wrapped in braces. The beginning of a synchroniser program is indicated by the keyword *synch*.

The header contains the synchroniser's name and the channel signature. The name is an ASCII string that follows the C convention.

The body lists the state and store variable declarations and the states of the underlying finite state machine. Each state is supplemented by a list of transitions. Each transition declares its triggering condition which includes an optional guarding state expression, an optional list of actions, and, finally, the destination state.

3.2.0.8 Channel signature

The channel signature defines the input and output channels of the synchroniser and their bracketing depths. The synchroniser header (Fig. 3.3) declares the synchroniser *min* with two input channels that are connected to the ports *a*, *b* and two output channels that are connected to the ports *c*, *d*. If the bracketing depths of the channels are not specified, they are assumed to be 0. Thus, the bracketing depth of the channel *a* is 0.

The input channel depth -1 indicates that the input channel is ignored in the synchroniser program. The output channel with the depth -1 must not have data sent to them.

```
synch min (a, b:p | c:2, d:p+1)
```

FIGURE 3.3: The synchroniser header

The *AstraKahn* synchroniser allows one to declare constant and configurable integer depths for the input and output channels. In addition, the depth of an output channel can be specified with an integer shift to the configurable input channel depth.

The input channels are required to have the bracketing depths specified in the signature. Thus, the channel *a* (i.e. the channel connected to the port *a*) of *min* must have zero bracketing depth. The channel *b* has a configurable bracketing depth *p*. The actual values of configurable bracketing depths of input channels are determined by the *AstraKahn* compiler by tracing bracketing depth relations over the application network.

The output channels of a synchroniser are guaranteed to have the bracketing depths specified in its channel signature. Thus, the synchroniser *min* must send messages to the output channel *c* at depth 2. The output channel *d* must have the bracketing depth $p + 1$, i.e. the depth of the input channel *b* plus 1.

3.2.0.9 Variable declaration

The start of a state variables declaration is marked by the keyword *state*. A state variable may be either an unsigned integer of some size or a C-style enumeration. State and store variable names are user-defined identifiers. A user-defined identifier is an ASCII string that follows the C convention.

Line 1 in Fig. 3.4 declares state variables *a*, *b*, *c* of size 4. Thus, all three variables are declared to have integer values in the range $[0; 15]$. Generally, a state variable of size *n*

has integer values in the range $[0; 2^n - 1]$. State variable c is initialised with an integer value 0.

The state variable foo that is declared on the line 2 in Fig. 3.4 can only be assigned the values d , e and f specified in the enumeration. The enumeration values are integer constants. If they are not specified explicitly, they are assigned consecutive positive integers starting with 0. Thus, the variable foo has integer values $d = 0$, $e = 1$ and $f = 2$.

The values can be specified explicitly (see line 3 in Fig. 3.4)

Integer state variables and enum values can be mixed freely in state expressions. Enum values are interpreted as integer constants.

```

1 state int(4) a, b, c = 0;
2 state enum(d, e, f) foo;
3 state enum(x = 1, y = 2, z = 4) bar;
4 store msg_a, msg_b;
```

FIGURE 3.4: State and store variables declaration

A store variable declaration begins with the keyword *store*. Line 4 in Fig. 3.4 declares state variables msg_a and msg_b . Store variables do not need an explicit type specification; their types are determined on the first assignment to the variable.

All the state and store variables are global to all the states.

3.2.0.10 States and transitions

States and transitions of the synchroniser define which channels are read and in what order. Fig. 3.5 presents the code of the binary zip synchroniser's state machine. Line 1 declares the start state of the synchroniser. The *on* clause indicates the beginning of the transition list. In the start state the zip2 synchroniser accepts messages from both input channels a and b . State and store expressions associated with the transition and the destination state are specified in the braces.

When the zip2 synchroniser is in the start state and it receives a message from channel a , the underlying state machine transitions to state $s1$. In this state the synchroniser can only receive messages from channel b since there's no transition triggered by channel a and defined in this state. When the message on channel b is received, the state machine transitions to the start state. Lines 10-13 define similar behaviour in state $s2$.

The synchroniser language supports top-down prioritised transition scopes. They are marked by the *elseon* keyword. A synchroniser in state foo in Fig. 3.6 accepts messages

```

1  start {
2    on:
3      a { goto s1;   }
4      b { goto s2;   }
5  }
6  s1 {
7    on:
8      b { goto start; }
9  }
10 s2 {
11   on:
12     a { goto start; }
13 }

```

FIGURE 3.5: State machine of the zip2 synchroniser

from channels connected to the ports *a*, *b*, *c* and *d*. When no destination state is specified for a transition, a synchroniser transitions to the current state. If all channels are ready at the same time in state *foo*, the synchroniser processes messages from either channel *a* or *b* first. When all messages from channels *a* and *b* are processed the synchroniser receives messages from channel *c*. If there're no messages in channels *a*, *b* and *c* the synchroniser receives messages from channel *d*.

```

1  foo {
2    on:
3      a { }
4      b { }
5    elseon:
6      c { }
7    elseon:
8      d { }
9  }

```

FIGURE 3.6: Prioritised transition scopes

3.2.0.11 State expressions

A state expression is a combination of integer constants, state variables and operators, which defines an integer value. The interpretation of a state expression follows the C rules of precedence and association. State expressions can be assigned to state variables. Under the assumption that the output channel has the infinite capacity a synthetic example in Fig. 3.7 counts the number of messages received from channel *a* between the arrivals of messages in channel *b*. Line 1 declares the 8-bit integer *count* and initialises it with 0. When a message from *a* is received the value of *count* increases by 1 (see line 5).

When a message from channel *b* is received the value of *count* is stored in the temporary variable *n*, set to 0 and then *n* is sent to the output channel.

```

1 state int(8) count = 0;
2 foo {
3   on:
4     a {
5       set count = [count + 1];
6     }
7   elseon:
8     b {
9       set n = [count], count = [0];
10      send count:[n] => c;
11    }
12 }

```

FIGURE 3.7: Use of state variables and expressions

The variable n does not have to be declared and is considered an alias for the integer expression. Temporary variables are available until the state machine of a synchroniser makes the next transition.

3.2.0.12 Triggering of a transition

The channel name on its own stands for the availability predicate for the corresponding channel, i.e. the condition that a message of any kind is available. Whether a transition takes place depends on the channel status and optionally the content of the message.

When a message is received on a channel, it can be matched with a pattern in order to extract parameters needed to select a specific transition. Line 3 of Fig. 3.8 checks whether a message received from the input channel a is a unique choice. Then the one and only variant is checked on whether it contains the label x . If it does, the value of x is stored in a temporary variable x . The tail of the message, i.e. all label-value pairs except for the value labeled x , is stored in a temporary variable t . Both x and t are available until the state machine makes the next transition.

```

1 foo {
2   on:
3     a.(x || t) { }
4     a.?v      { }
5     a.?v(x, y) { }
6     a.@[k]    { }
7 }

```

FIGURE 3.8: Message content extraction

To support message formats where several variants of a message are possible, a qualifier $?α$ is available as an input condition. It qualifies input messages as belonging to the $α$ variant. Line 4 of Fig. 3.8 checks whether the message received from channel a contains the variant v . Line 5 checks whether the message that contains the variant v with only two fields labeled x and y .

A channel carries a stream that consists of messages and possibly segmentation marks. Line 6 in Fig. 3.8 checks if the message is a segmentation mark of the depth equal to k . The depth of a segmentation mark can be compared with a state expression.

Several different channels can be tested in any given state, however, once the readiness of a channel is established, the synchroniser is committed. Hence the set of conditions applied to the message on any input channel must be exhaustive. In Fig. 3.8 it is not, because there is no pattern for messages that do not contain the field label x , the variant v and are not a segmentation mark of depth k at the same time. In this case the final clause *a.else*; is assumed. This clause discards the input message and transitions the synchroniser back to its current state.

A transition can be guarded by a state expression. In this case the transition fires only if the guarding expression evaluates to true. The synchroniser in Fig. 3.9 sends every 256-th message to the output channel. Line 1 declares the 8-bit state variable i and initialises it with 0. The variable is incremented every time a message from channel a is received, except when it reaches 255, in which case it is reset to 0 and the received message is sent down channel c .

Values that are matched from the message can be used in guarding state expressions.

```

1 state int(8) i;
2 start {
3   on:
4     a & [i < 255] {
5       set i = [i + 1];
6     }
7     a & [i = 255] {
8       set i = [0];
9       send this => c;
10    }
11 }
```

FIGURE 3.9: Use of guarding state expressions

3.2.0.13 Store expressions and sending messages

Store expression is a mechanism to combine data. In *AstraKahn* data are typed. Types are CAL terms. The result of the store expression can be either saved in a store variable or sent down an output channel.

The example in Fig. 3.10 demonstrates the use of store expressions and the *send* clause. In the start state the synchroniser receives messages from channel a that has the label n in it. Line 5 increments the value under the label n and stores it in the store variable ma under the label n together with the tail t . The operator $'$ applied to the variable x

creates the record ' $x' : value(x)$ '. This is a shorthand useful for the avoidance of tedious notation.

```

1 store ma;
2 start {
3   on:
4     a.(n || t) {
5       set ma = (n:[n+1] || t);
6       goto s1;
7     }
8 }
9 s1 {
10  on:
11    b {
12      send ma || this => c;
13      goto start;
14    }
15 }
```

FIGURE 3.10: Use of store expressions and the *send* clause

A message received on a channel is referred to by the keyword *this* within the active transition. In state *s1* the synchroniser receives messages from channel *b*. When a message is received, it is concatenated with the store variable *ma* (see line 12) and sent to the output channel *c*.

3.3 Execution Order of Synchroniser

In order to achieve the lowest latency, an *AstraKahn* synchroniser exploits a non-deterministic behaviour. In a certain state more than one input channel may be ready, however, a state machine receives one input message at a time. The synchroniser does not take any transition that potentially causes sending to a blocked channel. Of the transitions that do not send to blocked channels, which one will be triggered is defined by the fairness policy: when more than one transition is possible in a given state, all choices will be made with the same frequency. If transitions are prioritised, the choices are made within each scope at first.

Once the transition has been taken and the associated actions have been executed, the synchroniser transitions to the destination state. In order to avoid transitioning to a state in which there are no ready input channels the synchroniser may be provided with a set of destinations to choose. In the synchroniser language this is expressed as a goto-clause with multiple destinations. The algorithm in Fig. 3.11 defines the ordering of the choices that a synchroniser makes during the execution.

Require: The current state of the synchroniser (*curr_state*)
Ensure: The state, to which the synchroniser transits from the current state

```

1: function RUN(curr_state)
2:   ReadyInputs  $\leftarrow$  ready channels that are read in curr_state
3:   for each channel in ReadyInputs do
4:     trans  $\leftarrow$  transitions from curr_state that read from channel
5:     if  $\exists t \in \text{trans} \wedge t \text{ causes sending to a blocked channel}$  then
6:       remove channel from ReadyInputs
7:     end if
8:   end for
9:   if ReadyInputs =  $\emptyset$  then
10:    return curr_state
11:  end if

12:  channel  $\leftarrow$  the least frequently taken channel from ReadyInputs
13:  message  $\leftarrow$  fetch a message from channel
14:  for each priority P in curr_state do  $\triangleright$  iterate over transition priorities from
    the highest to the lowest
15:    trans  $\leftarrow$  all transitions that read from channel with the priority P
16:    valid_trans  $\leftarrow$  all transitions from trans with satisfied conditions
17:    else_trans  $\leftarrow$  the .else transition from trans with the satisfied condition
18:    if valid_trans =  $\emptyset$  then
19:      if else_trans = nil then
20:        continue
21:      else
22:        choose(else_trans)
23:      end if
24:    else
25:      choose(the least frequently taken transition from valid_trans)
26:    end if
27:  end for
28:  if no transition has been chosen then
29:    return curr_state
30:  end if
31:  act()  $\triangleright$  perform the actions associated with the chosen transition
32:  ImmediateStates  $\leftarrow$  states listed by the goto clause

33:  for each state in ImmediateStates do
34:    ReadyInputs  $\leftarrow$  ready channels that are read in state
35:    trans  $\leftarrow$  transitions from state that read from ReadyInputs
36:    if ReadyInputs =  $\emptyset \vee \exists t \in \text{trans} \wedge t \text{ causes sending to a blocked channel}$ 
then
37:      remove state from ImmediateStates
38:    end if
39:  end for
40:  if ImmediateStates  $\neq \emptyset$  then
41:    return the least frequently taken state from ImmediateStates
42:  end if
43:  return the least frequently taken state from the goto clause list
44: end function

```

FIGURE 3.11: The execution of a synchroniser

3.4 The Implementation of the *aksync* Compiler

In this section we describe the implementation of the *AstraKahn* synchroniser compiler *aksync*. At the current stage of the *AstraKahn* software stack development the *aksync* compiler is highly integrated into the *AstraKahn* runtime system prototype. Similar to the runtime system the *aksync* compiler is implemented in Python. It generates an intermediate representation of the synchroniser program which is passed to the synchroniser runtime.

The lexical and syntax analysers are implemented using PLY [14] - an implementation of lex and yacc for Python. The semantic analyser performs semantic and type checking. The code generator emits the synchroniser's runtime data structure and derives its CAL passport.

3.4.1 Lexical Analysis

3.4.1.1 Lexical analyser

The lexical analyser reads the stream of characters making up the source program and groups the characters into lexemes. For each lexeme, the lexical analyser produces a token of the form $\langle name, value \rangle$, which it passes to the syntax analyser. For tokens that do not need the *value*, such as punctuation, reserved words and keywords, the second component is omitted. Those are given in Appendix A.2.

PLY implements the way in which traditional tools work. Specifically, the Python lex provides an external interface in the form of a `token()` function that returns the next valid token on the input stream. Token positional information, which is useful in the context of error handling, is managed by the Python yacc.

3.4.1.2 Preprocessor

The original synchroniser language in [1] provides integer configuration parameters to avoid having to trivially alter synchroniser programs. They are specified in brackets between the synchroniser's name and its channel signature. A program that has to be compiled with uninstantiated input parameters potentially makes the analyses in the compiler more conservative. Thus, all parameters should be instantiated before compilation starts.

For the sake of simplicity, we implement substitution for the free parameters in a tiny lexical preprocessor. The lexical preprocessor requires only lexical analysis. It substitutes tokenized character sequences for other tokenized character sequences according to some user-defined rules. The preprocessor that has been implemented does not support any directives and only performs macro substitution. The compiler reads macros from its invocation command and then passes them to the preprocessor.

With the above implementation, configuration parameters do not have to be specified in the synchroniser program. However, such an implementation of configuration parameters has a serious drawback: a macro defines a blind substitution. There is no way to make sure that it defines a rule for the substitution of an integer parameter because lexical analysis does not know anything about program structure and semantics. However, at least rules that are obviously not suitable for integer parameter substitution can be diagnosed.

3.4.2 Syntax Analysis

3.4.2.1 Syntax alanyser

The syntax analyser obtains a string of tokens from the lexical analyser and verifies that it can be generated by the grammar of the source language. The syntax analyser is expected to report syntax errors in an intelligible way and recover from common errors in order to continue processing of the remainder of the program. For well-formed programs, the syntax analyser constructs an intermediate representation of the program and passes it to the rest of the compiler for further processing. The *aksync* implementation is based on the abstract syntax tree representation.

The Python yacc generates a syntax-directed translator. Syntax-directed translation is done by attaching program fragments to productions in a grammar. The program fragment or so-called semantic action is executed when the production is used during syntax analysis. The combined result of these executions produces the intermediate representation of a program in the order induced by the syntax analysis. The Python yacc accepts source language syntax specification in context-free grammar form. The grammar for the synchroniser language is given in [Appendix A.1](#).

The syntax error handling and recovery mechanism of PLY is similar to the one of Unix's yacc. During syntax analysis when a syntax error is detected the analyser switches to recovery mode in order to continue futher analysis to detect the remaining errors. In our implementation we do not use recovery mode. Instead, we focus on reporting an error in the best possible way. We augment the grammar for the language with productions that

generate erroneous constructs. The Python yacc provides a special token `error` that acts as a wildcard for any erroneous input. An analyser constructed from a grammar augmented by these error productions is useful for detecting anticipated errors.

3.4.2.2 Symbol table

Symbol tables are data structures used by compilers to hold information about identifiers coalesced from the program's source code. A semantic action puts information about identifier x into the symbol table, when the declaration of x is analysed, and uses it when necessary.

The scope of a declaration is the portion of the program to which the declaration applies. In synchroniser code state and store variables are visible for all states and transitions. State expression aliases and pattern-matched variables are visible only within the current transition. We have implemented scopes using a chained symbol table approach described in [15]. We set up a root symbol table for state and store variables and separate symbol tables for each transition. The tables are chained as illustrated in Fig. 3.12.




FIGURE 3.12: Scoping structure of a synchroniser program using the zip2 synchroniser as an example

The symbol table implementation supports three operations: create new symbol table, put new entry in table and get entry for identifier from table. In the sequel, we refer to these operations as *NewSymtab*(symtab), *symtab.put*(ID, value) and *symtab.get*(ID). The pseudo-code implementations of *symtab.get* and *symtab.put* are given in Fig. 3.13 and Fig. 3.14 respectively.

The synchroniser language does not permit using of the reserved words (Appendix A.2) as identifiers. This is checked before putting an identifier in the symbol table.

3.4.3 Static Analysis

Static analysis includes:

```

1: function GET(symtab, id)
2:   while symtab is not nil do
3:     tmp  $\leftarrow$  get(symtab, id)
4:     if tmp then
5:       return tmp
6:     else
7:       symtab  $\leftarrow$  previous(symtab)  $\triangleright$  previous(symtab) returns a symbol table
        of the most-closely outer scope
8:     end if
9:   end while
10:  return not_found
11: end function

```

FIGURE 3.13: Getting an entry for an identifier from the chained symbol table

```

1: function PUT(symtab, id, value)
2:   if id is reserved then
3:     error
4:   end if
5:   tmp  $\leftarrow$  get(symtab, id)
6:   if not tmp then
7:     symtab  $\leftarrow$  (id, value)
8:   else
9:     error
10:  end if
11: end function

```

FIGURE 3.14: Putting a new entry in the symbol table

- Semantic checking. Constraints such as an identifier is declared at most once in a scope
- Type checking. The type rules of a language assure that an operator or function is applied to the right number and type of operands.

3.4.3.1 Semantic checking

In this section we describe the semantic checks that are not enforced by the grammar.

The synchroniser language requires identifiers except for state expression aliases to be declared before they are used. Moreover, an identifier must be declared at most once in a scope. The channels, on which the transitions in the synchroniser are made, must be declared in the channel signature as well as the channels where messages are sent. In order to check those, we maintain three symbol tables for identifiers, input channels and output channels. The scheme given in Fig. 3.1 shows how the symbol tables are

managed and used during the syntax analysis. The symbol tables are initialised before the syntax analyser runs, as shown in Fig. 3.15.

The attribute *type* is added to each entry in the symbol table. State variables are assigned type *integer*. Non-integer variables, whose structure is unknown, are assigned a special type *void*, which stands for a variable CAL term. A detailed information about the synchroniser language type system is provided in section 3.4.3.2.

```

function INIT
  InChantab  $\leftarrow$  NewSymtab(nil)
  OutChantab  $\leftarrow$  NewSymtab(nil)
  RootSymtab  $\leftarrow$  NewSymtab(nil)
end function

```

FIGURE 3.15: Initialisation of symbol tables

The synchroniser compiler checks if the transition diagram of a synchroniser is connected. The algorithm given in Fig. 3.16 walks the abstract syntax tree *synch* and constructs two sets: the set of the synchroniser state labels *StateSet* and the set of the goto labels *GotoSet*. The set $GotoSet \setminus StateSet$ contains goto labels that point to non-existent states. If this set is not empty the compilation reporting an error. The set $StateSet \setminus (GotoSet \cup \text{'start'})$ contains unreachable states that are eliminated, also reporting an error. The algorithm also checks if the labels of the synchroniser states are unique.

A synchroniser program must have the state labeled 'start'. The existence of this state is checked once the *StateSet* is constructed.

The synchroniser language provides two types of state variable: an integer of the specified size and an enumeration. The size defines the value range of the integer. The value range of the enumeration is specified in the variable declaration. Fig. 3.17 gives formulae for the value range computation.

For a state expression that evaluates to an integer we can check if the computed value belongs to the assignment-destination value-range. The symbol table stores the value-range information for integer entries and provides an interface to it in the form of boolean function `check_range(id, value)`. The function returns *true* if the *value* fits in the *id* value range and *false* otherwise. Fig. 3.19 shows how the check is integrated into the syntax analyser.

Tests We have developed a test suite for the semantic analyser. It uses the standard python unit testing framework *unittest*. The tests expand the abstract syntax tree into a nested list and compare the result with the expected value.

Production	Semantic Action
$\langle \text{synch} \rangle ::= \text{'synch'} \langle ID \rangle \text{'('} \langle \text{input} \rangle \text{'}, \text{' } \langle \text{input} \rangle \text{'})^* \text{'}, \text{' } \langle \text{output} \rangle \text{'}, \text{' } \langle \text{output} \rangle \text{'})^* \text{'}, \text{' } \{ \text{' } \langle \text{decl} \rangle^* \langle \text{state} \rangle^+ \text{'}, \text{' } \}$	$\text{ /* synch is the abstract syntax tree of the source code. */ }$
$\langle \text{input} \rangle ::= \langle \text{chan} \rangle \text{'('} \text{'}, \text{' } \langle ID \rangle \text{' } \langle \text{NUMBER} \rangle \text{'})'$ $\langle \text{chan} \rangle ::= \langle ID \rangle$	$\text{foreach } \langle \text{chan} \rangle$ $\text{InChantab.put}(\langle \text{chan} \rangle, \text{type}=\text{void})$
$\langle \text{output} \rangle ::= \langle \text{chan} \rangle \text{'('} \text{'}, \text{' } \langle \text{depth_exp} \rangle \text{'})'$ $\langle \text{chan} \rangle ::= \langle ID \rangle$	$\text{foreach } \langle \text{chan} \rangle$ $\text{OutChantab.put}(\langle \text{chan} \rangle, \text{type}=\text{void})$
$\langle \text{decl} \rangle ::= \text{'store'} \langle \text{store_id_list} \rangle \text{'}, \text{'}$ $\quad \text{'state'} \langle \text{type} \rangle \langle \text{state_id_list} \rangle \text{'}, \text{'}$ $\langle \text{store_id_list} \rangle ::= \langle \text{id_list} \rangle$ $\langle \text{state_id_list} \rangle ::= \langle \text{id_list} \rangle$	$\text{foreach id in } \langle \text{store_id_list} \rangle$ $\text{Symtab.put}(\text{id}, \text{type}=\text{void})$ $\text{foreach id in } \langle \text{state_id_list} \rangle$ $\text{Symtab.put}(\text{id}, \text{type}=\text{int})$
$\langle \text{trans_stmt} \rangle ::=$ $\langle \text{trans_name} \rangle \text{'('} \text{'}, \text{' } \langle \text{condition} \rangle \text{' } \text{'\&' } \langle \text{int_exp} \rangle \text{' } \langle \text{actions} \rangle$	$\text{if not InChantab.get}(\langle ID \rangle)$ error
$\langle \text{trans_name} \rangle ::= \langle ID \rangle$	$\text{Symtab} = \text{NewSymtab}(\text{RootSymtab})$ $\text{Symtab.put}(\text{'this'}, \text{type}=\text{void})$
$\langle \text{condition} \rangle ::= \text{'@'} \langle \text{segmark} \rangle$ $\quad \text{'?' } \langle ID \rangle$ $\quad \text{'else'}$ $\langle \text{segmark} \rangle ::= \langle ID \rangle$	$\text{if not Symtab.get}(\langle \text{segmark} \rangle)$ error
$\langle \text{condition} \rangle ::= \text{'?' } \langle ID \rangle \text{'('} \text{'}, \text{' } \langle \text{id_list} \rangle \text{' } \text{' ' } \langle \text{tail} \rangle \text{' } \text{'})'$	$\text{foreach id in } \langle \text{id_list} \rangle \cap \langle \text{tail} \rangle$ $\text{Symtab.put}(\text{id}, \text{type}=\text{void})$ $\text{tmp} = \text{Symtab.get}(\text{'this'})$ $\text{tmp.type} = \{ \text{'p}_1': \text{void}, \dots, \text{'p}_n': \text{void} \}, \text{ where } p_1,$ $\dots p_n \text{ are elements of } \langle \text{id_list} \rangle$ $\text{if } \langle \text{tail} \rangle$ $\text{tmp.type} = \text{tmp.type} \cup \langle \text{tail} \rangle$ $\text{tmp} = \text{OutChantab.get}(\langle ID \rangle)$ if not tmp error
$\langle \text{send_stmt} \rangle ::= \text{'send'} \langle \text{dispatch} \rangle \text{'('} \text{'}, \text{' } \langle \text{dispatch} \rangle^* \text{'}, \text{'}$ $\langle \text{dispatch} \rangle ::= \langle \text{msg_exp} \rangle \text{'=>' } \langle ID \rangle$	$\text{tmp} = \text{OutChantab.get}(\langle ID \rangle)$ if not tmp error

TABLE 3.1: Symbol tables management and duplicate declaration checking scheme

```

1: function GET_STATES(synch)
2:   StateSet  $\leftarrow \square$ 
3:   GotoSet  $\leftarrow \square$ 
4:   for each state in state_list(sync) do
5:     if label(state)  $\in$  StateSet then
6:       error ▷ The state label is not unique
7:     else
8:       s  $\leftarrow s \dots \text{label}(\text{state})$ 
9:     end if
10:    for each trans in trans_list(state) do
11:      for each gotostate in goto_list(trans) do
12:        if gotostate  $\notin$  GotoSet then
13:          GotoSet  $\leftarrow$  GotoSet  $\dots$  gotostate
14:        end if
15:      end for
16:    end for
17:  end for
18:  return (StateSet, GotoSet)
19: end function

```

FIGURE 3.16: Construction of *StateSet* and *GotoSet*

Type	Values
$\text{int}(n)$	$[0; 2^n - 1] \cap \mathbb{Z}$
$\text{enum}(a_1, a_2, \dots, a_n)$	$[0; n - 1] \cap \mathbb{Z}$
$\text{enum}(a_1 = N_1, a_2 = N_2, \dots, a_n = N_n)$	N_1, N_2, \dots, N_n

FIGURE 3.17: Computing the value range of an integer variable

3.4.3.2 Type checking

The design of the type checker is based on information about the syntactic constructs in the language, the notion of types and the rules for assigning types to language constructs. The type of a language construct is denoted by a type expression. Informally, a type expression is either a basic type or the application of an operator called a type constructor to other type expressions. A collection of rules for assigning type expressions to language constructs is called a type system.

The communication protocol of the *AstraKahn* runtime system prototype supports only choices. When a synchroniser reads a message from its input channel, the record that belongs to the variant specified in the synchroniser transition is instantiated³. We take this into account in implementing the type system of the synchroniser language. Because the synchronisers can read values only of the fields that are known to be integer, the basic types in the synchroniser language are integer and CAL variable. Integer is the type of state variable. CAL variables are building blocks for the only constructed type

³this means that store variables cannot keep multiple variants

```

1: function  $||(r_1, r_2)$ 
2:   if  $len(r_1) \leq len(r_2)$  then
3:      $r_{iter} \leftarrow r_1$ 
4:      $r \leftarrow r_2$ 
5:   else
6:      $r_{iter} \leftarrow r_2$ 
7:      $r \leftarrow r_1$ 
8:   end if  $\triangleright r$  is the record that contains fewer label-value pairs
9:   for each label-value pair  $(l, v)$  in  $r_{iter}$  do
10:    if  $r(l)$  then  $\triangleright$  if label  $l$  exists in  $r$ 
11:       $r(l) \leftarrow union(r(l), v)$ 
12:    else
13:       $r(l) \leftarrow v$ 
14:    end if
15:  end for
16:  return  $r$ 
17: end function

```

FIGURE 3.18: The record type constructor $||$

in the synchroniser language – a record. A record is constructed by the concatenation of two records. The pseudo-code of the record constructor $||$ is given in Fig. 3.18. The record constructor is obviously commutative and associative.

The case when a label-value pair labeled l exists in both operand records r_1 and r_2 is indicated with $union(r_1(l), r_2(l))$. The CAL solver resolves which option to take during the constraint aggregation pass in the *AstraKahn* compiler.

We describe the type systems in terms of grammar productions and corresponding semantic actions. The type systems for state and store expressions is given in Fig. 3.20 and Fig. 3.21 respectively. The synthesized attribute *type* for an expression $\langle E \rangle$ gives the type of the expression assigned by the type system for the expression generated by $\langle E \rangle$. The type system for statements is given in Fig. 3.19. It assures that the left hand side can be assigned to.

3.4.4 Code Generation

3.4.4.1 Synchroniser runtime code

The compiler generates a data structure that is passed to the *AstraKahn* runtime system for the interpretation.

Production	Semantic Action
$\langle assign \rangle ::= \langle dest \rangle \text{ '=' } [\langle int_exp_c \rangle]$ $\langle dest \rangle ::= \langle ID \rangle$	$tmp = \text{Symtab.get}(\langle dest \rangle)$ <i>if not</i> tmp $\text{Symtab.put}(\langle dest \rangle, type=int)$ <i>else</i> <i>if</i> tmp.type $\neq int$ <i>error</i> <i>if</i> $\langle int_exp_c \rangle$ evaluates to int <i>if not</i> check_range($\langle dest \rangle$, eval($\langle int_exp_c \rangle$)) <i>error</i>
$\langle assign \rangle ::= \langle dest \rangle \text{ '=' } \langle data_exp \rangle$ $\langle data_exp \rangle ::=$ $(\langle data \rangle \mid \text{'(' } \langle data \rangle \text{'})$	$tmp = \text{Symtab.get}(\langle dest \rangle)$ <i>if not</i> tmp <i>error</i> <i>if</i> tmp.type == int <i>error</i> $tmp.type = \langle data_exp \rangle.type$

FIGURE 3.19: Type system for statements

Production	Semantic Action
$\langle int_exp_c \rangle ::= \langle NUMBER \rangle \mid \langle ID \rangle$	$tmp = \text{Symtab.get}(\langle ID \rangle)$ <i>if not</i> tmp <i>error</i> <i>if</i> tmp.type $\neq int$ <i>error</i> $\langle int_exp_c \rangle.type = int$
$\langle int_exp_c \rangle ::= \text{'(' } \langle int_exp_c_1 \rangle \text{'}'$ $\mid \text{'-'} \langle int_exp_c_1 \rangle$ $\mid \text{'!'} \langle int_exp_c_1 \rangle$	$\langle int_exp_c \rangle.type = \langle int_exp_c_1 \rangle.type$
$\langle int_exp_c \rangle ::=$ $\langle int_exp_c_1 \rangle \langle op \rangle \langle int_exp_c_2 \rangle$ $\langle op \rangle ::= \text{'+'} \mid \text{'-'} \mid \text{'*'} \mid \text{'/'} \mid \text{'%'}$ $\mid \text{'<<'} \mid \text{'>>'} \mid \text{' '} \mid \text{'\&'} \mid \text{'^'}$ $\mid \text{'<'} \mid \text{'>'} \mid \text{'=='} \mid \text{'!='} \mid \text{'<='}$ $\mid \text{'>='} \mid \text{'\&\&'} \mid \text{' '}$	<i>if</i> $\langle int_exp_c_1 \rangle.type == int$ <i>and</i> $\langle int_exp_c_2 \rangle.type == int$ $\langle int_exp_c \rangle.type = int$ <i>else</i> <i>error</i>

FIGURE 3.20: Type system for state expressions

Production	Semantic Action
$\langle data \rangle ::= \langle item_list \rangle$ $\langle item_list \rangle ::= \langle item \rangle [' ' \langle item \rangle]^*$	<i>foreach</i> item <i>in</i> $\langle item_list \rangle$ $\langle data \rangle.type = \langle data \rangle.type \mathbf{item}.type$
$\langle item \rangle ::= \mathbf{'this'}$	$\mathbf{tmp} = \text{Symtab.get}(\mathbf{'this'})$ $\langle item \rangle.type = \mathbf{tmp.type}$
$\langle item \rangle ::= \langle ID \rangle$	$\mathbf{tmp} = \text{Symtab.get}(\langle ID \rangle)$ <i>if not</i> \mathbf{tmp} <i>error</i> <i>if</i> $\mathbf{tmp.type} == \mathbf{int}$ <i>error</i> $\langle item \rangle.type = \mathbf{tmp.type}$
$\langle item \rangle ::= \mathbf{' ' \langle ID \rangle}$	$\mathbf{tmp} = \text{Symtab.get}(\langle ID \rangle)$ <i>if not</i> \mathbf{tmp} <i>error</i> $\langle item \rangle.type = \{ \mathbf{'ID'} : \mathbf{tmp.type} \}$
$\langle item \rangle ::= \langle ID \rangle \mathbf{' : ' \langle rhs \rangle}$ $\langle rhs \rangle ::= \langle int_exp \rangle \langle rhs_ID \rangle$ $\langle rhs_ID \rangle ::= \langle ID \rangle$	<i>if</i> $\langle rhs_ID \rangle$ $\mathbf{tmp} = \text{Symtab.get}(\langle rhs_ID \rangle)$ <i>if not</i> \mathbf{tmp} <i>error</i> $\mathbf{tmp} = \text{Symtab.get}(\langle ID \rangle)$ $\langle item \rangle.type = \{ \mathbf{'ID'} : \langle rhs \rangle.type \}$

FIGURE 3.21: Type system for store expressions

3.4.4.2 Synchroniser passport

Synchronisers rely on both TPL and CAL for their definition. The CAL aspects of a synchroniser are confined to the CAL terms for its input and output channels. Those terms are not straightforward since they have to be fairly generic to match the broadest possible formats of producer and consumer messages involved in the act of synchronisation. On the other hand, the synchroniser passport is produced solely on the basis of the synchroniser code, exclusively by program analysis; the programmer does not supply an explicit passport for this.

The first thing that requires CAL is the input interface. As mentioned above, the communication protocol of the *AstraKahn* runtime system prototype supports only choices. The use of a variant on a channel in any transition on that channel immediately associates the *choice* term structure with it. For example, a channel *a* that is tested on variants $?v$ and $?w$ in transitions has a term comparable with $(\mathbf{: 'v':\$vterm, 'w':\$wterm} || \$tail \mathbf{:})$, where the variables $\$vterm$ and $\$wterm$ represent the terms for variants *v* and *w* and $\$tail$ represents the choice term that contains the rest of the variants. If a choice is known to carry a single variant, the variant is labeled *uniq*. A transition that does not specify a variant label expects a unique message variant.


```

1: function INPUT_TERM(synch)
2:   CondDict  $\leftarrow$  nil
3:   for each state in state_list(sync) do
4:     for each trans in trans_list(state) do
5:       cond  $\leftarrow$  get_condition(trans)
6:       if port  $\in$  CondDict then
7:         CondDict(port)  $\leftarrow$  CondDict(port) .. cond
8:       else
9:         CondDict  $\leftarrow$  CondDict .. (port, cond)
10:      end if
11:    end for
12:  end for
13:  InputTerm  $\leftarrow$  nil
14:  for each (port, cond_set) in CondDict do
15:    PortTerm  $\leftarrow$  nil
16:    for each cond in cond_set do
17:      variant  $\leftarrow$  get_variant(cond)
18:      this  $\leftarrow$  Symtab.get('this')
19:      if variant is unique then
20:        if PortTerm  $\neq$  nil then error
21:        end if
22:        PortTerm  $\leftarrow$  (: 'uniq':this :)
23:        break
24:      else
25:        PortTerm  $\leftarrow$  PortTerm || this
26:      end if
27:    end for
28:    InputTerm  $\leftarrow$  InputTerm .. (port, PortTerm)
29:  end for
30:  return InputTerm || $tail
31: end function

```

FIGURE 3.22: Construction of the synchroniser input term

The symbol table for a transition maintains a special entry ‘this’. It holds the term of the message accepted by the transition. The algorithm in Fig. 3.22 walks the synchroniser transitions and constructs the CAL input term for the synchroniser⁴.

Since a choice is in fact a collection of label-record pairs, the constructor || (Fig. 3.18) can be applied for choices. When *InputTerm* contains a label-value pair (‘v’, *value*) and the label ‘v’ occurs in another transition, the union || of these two choices results in the following:

$$\begin{aligned}
& \text{union}((: \text{‘v’} : \$vterm_1 \mid \$tail_1 :), (: \text{‘v’} : \$vterm_2 \mid \$tail_2 :)) \\
& = \quad (: \text{‘v’} : \text{union}(\$vterm_1, \$vterm_2) \mid (\$tail_1 \mid \$tail_2) :)
\end{aligned}$$

⁴The algorithm relies on the abstract syntactic tree structure built by the compiler (see Appendix. A.3)

```

1: function OUTPUT_TERM(synch)
2:   DispatchDict  $\leftarrow$  nil
3:   for each state in state_list(sync) do
4:     for each trans in trans_list(state) do
5:       send  $\leftarrow$  get_send(trans)
6:       (msg, port)  $\leftarrow$  (get_msg(send), get_port(send)
7:       if port  $\in$  DispatchDict then
8:         DispatchDict(port)  $\leftarrow$  DispatchDict(port) .. msg
9:       else
10:        DispatchDict  $\leftarrow$  DispatchDict .. (port, msg)
11:      end if
12:    end for
13:  end for
14:  OutputTerm  $\leftarrow$  nil
15:  for each (port, msg_set) in DispatchDict do
16:    PortTerm  $\leftarrow$  nil
17:    for each msg in msg_set do
18:      if msg is MsgData then  $\triangleright$  msg matches  $['?'\langle ID \rangle]\langle data \rangle$ 
19:        (variant, data)  $\leftarrow$  (get_variant(msg), get_data(msg))
20:        for each item in data do
21:          if item is ItemVar or item is ItemPair then  $\triangleright$  item is either 'id'
22:            or id.value
23:            (lhs, rhs)  $\leftarrow$  expand(item)
24:            PortTerm  $\leftarrow$  PortTerm || {lhs.type(rhs)}
25:          else  $\triangleright$  item is either this or id
26:            PortTerm  $\leftarrow$  PortTerm || type(id)
27:          end if
28:        end for
29:      end if
30:    OutputTerm  $\leftarrow$  OutputTerm .. (port, PortTerm)
31:  end for
32:  return OutputTerm || $tail
33: end function

```

FIGURE 3.23: Construction of the synchroniser output term

Now consider the output interface. The algorithm in Fig. 3.23 collects the dispatches for every output channel and combines them into the output term of every output channel.

The function *type*(*id*) performs the symbol table lookup and returns the *type* attribute value for the 'id' entry.

3.5 Discussion and Future Work

In this chapter we have presented the language for *AstraKahn* synchronisers. We have developed the language compiler that generates the CAL passport of the synchroniser. The

message format in *AstraKahn* is based on the Message Definition Language (MDL) with the restriction that data on streams are organised as collections of alternative records of label-value pairs. A value in a record can have any structure that is allowed by the MDL. The synchroniser that we have presented matches only a top-level structure of a message. The MDL generates a much broader set of terms than the current version synchroniser can synchronise; however, it needs to be elaborated whether the implementation of lower level synchronisation in synchronisers is useful for the real world applications.

The current version of the language does not define flow inheritance in synchronisers. The synchroniser *code* only needs to access the label-value parts of the message it matches. Thus, the flow inheritance should be supported outside of the synchroniser definition and the question how it should be done is left open.

The compiler we have implemented does not optimise the code deeply. We have only implemented an elimination of unused states. However, a broader set of dead code elimination optimisations can be implemented⁵.

⁵E.g. a conditional transition can be eliminated if its condition always evaluates to *false*

Chapter 4

Serial Replication in *AstraKahn*

In this chapter we explain the machinery behind the serial replication in *AstraKahn* and the role of synchronisers in it. We introduce the concept of the forward fixed point for the replication pipeline and show how it is used to organise the output from the infinite chain of replicas. In order to suppress the growth of the replica chain, we present the concept of the reverse fixed point and show how it is used to optimise some replicas at the head of the chain.

4.1 *AstraKahn* Approach to Serial Replication


The serial replication combinator creates a conceptually infinite number of copies of its operand network, and connects them in a chain. Replication is demand-driven: the replicas are created dynamically. A fresh replica is *inactive*¹, hence it does not necessarily require significant resources since *AstraKahn* boxes are stateless and since synchronisers require no resources in their start state². Indeed the cost of replication is only felt when the replicas are active, which is the case from the time that the first message is received until all messages have left the replica and all its synchronisers have returned to their start states.

In S-NET, the output from a replication pipeline is based on the record subtyping in the type system. The replication combinators in S-NET require the programmer to specify a termination pattern, so that each record that is a subtype of this pattern leaves the replication pipeline through the output stream.

¹More generally, we call a replica inactive when all of its synchronisers are in their start states, none of its channels has messages in them and no box is running

²When a synchroniser transitions back to the start state, it flushes its store variables

In *AstraKahn* the output from the replication pipeline is defined using the concept of fixed point. From the mathematical point of view a fixed point of a function is an element of the function's domain that it maps to itself. That is to say, a function $f(x) : X \rightarrow Y$ has a fixed point at $x_0 \in X$ if $f(x_0) = x_0$. The serial replication combinator implements the computation³ shown in Fig. 4.1. After the n -th replica has processed the message $f^{(n-1)}(x) \neq x_0$ the computation reaches the fixed point $f^{(n)}(x) = x_0$, and the message x_0 is sent on to the output channel of the serial replication network f^* .

FIGURE 4.1: The recursive computation in *AstraKahn*

The number of iterations n needed to reach the fixed point is not known in advance, meaning that in order to utilize the fixed point as shown in Fig. 4.1, *AstraKahn* must be able to detect the fixed point message right at the time it is produced by the n -th replica or later. Therefore, similar to S-NET, *AstraKahn* needs to be provided with a pattern that matches all the fixed point messages of the operand network. In S-NET, the serial replication combinator requires this pattern as one of its operands; by contrast, in *AstraKahn* the pattern is required to be deduced from the operand network by compiler analysis.

The chain of replicas grows as the computation progresses, however, in the example in Fig. 4.1 the computation is carried out only by a single replica in the tail of the chain. The replicas in the head of the chain have processed the message and are not used anymore. In order to suppress the growth of the chain, *AstraKahn* must detect such replicas and optimise the connection by removing them. Since *AstraKahn* boxes are stateless, an operand network can have a state that is fully defined by the states of its synchronisers. A replica of the operand network can be removed from the chain safely iff the replica is in a state, in which it forwards any message without change, i.e. any message it receives is its fixed point. We will call such a state of the replica a reverse fixed point state.

In the remainder of the chapter we will give formal definitions of a fixed point message and a reverse fixed point state, and will provide algorithms for the *AstraKahn* compiler to detect them. In the sequel, we will call a fixed point message a forward fixed point.

³ $f^{(n)}(x)$ denotes $\underbrace{f(f(\dots f(x)\dots))}_n$

4.2 Forward Fixed Point

Once the computation in a serial replication network has reached the fixed point, newly created replicas are known to transmit fixed point messages without change. *AstraKahn* does not analyse boxes⁴, so it can determine about the operand network behaviour only from its synchronisers. Thus, in order for the operand network to be analysable by *AstraKahn* it must contain a path that does not traverse boxes and which may traverse synchronisers. Because a newly created replica is inactive, and hence the synchronisers in it are in their start states, the start states of the synchronisers that belong to the path must have a special transition that sends the message on to the next synchroniser along the path. Since transitions can be conditional on the message content, the fixed point pattern, or rather the fixed point condition, can be present in these special transitions.

The existence of a forward fixed point requires the operand network to have some topological properties that are formally defined as follows. Consider a network N that has an input and an output channel, both named x .

Definition 1. The network N is said to have a forward fixed point in x if and only if the following requirements are satisfied:

1. There exists a condition $P(m)$ on the content of the message m received by the network on the input channel x under which it follows a unique non-branching path to the output channel x without traversing any boxes
2. The path⁵ can traverse synchronisers, but then whenever $P(m)$ is true and the synchroniser is in the start state, it must accept m and transition back to the start state while sending the message m on the path unchanged and without producing any other output

The condition P may not be unique for each network, and when it is not, the fixed point condition of the network is a disjunction of all such conditions. The condition can also be a tautology, in which case the forward fixed point is called unconditional. When the fixed point path traverses a single synchroniser, the fixed point condition is defined exclusively by the synchroniser transitions that loop around the start state and send on the accepted messages unchanged. When the path traverses several synchronisers, the fixed point condition of the network is a conjunction of the fixed point conditions of these synchronisers. We demonstrate the construction of the fixed point condition with the example operand network N depicted in Fig. 4.2.

⁴Except for the CAL passport generation

⁵In the sequel, we will call this path the fixed point path




FIGURE 4.2: Forming of the forward fixed point condition of a network

Apart from all the paths that traverse boxes, the operand network N has a unique path $[s_1, s_2, s_3]$ that traverses only synchronisers. The synchroniser s_2 has two transitions that loop around the start state and send on messages they accept unchanged with the firing conditions p_2^1 and p_2^2 . A message m is a fixed point for the synchroniser s_2 when it satisfies any of these conditions, i.e. $p_2^1(m) \vee p_2^2(m)$ is true. The synchronisers s_1 and s_3 have fixed point conditions p_1 and p_3 respectively. Then the fixed point condition of the network N is $p_1 \wedge (p_2^1 \vee p_2^2) \wedge p_3$.

Definition 1 requires the synchronisers that traverse the fixed point path to transition back to the start state after the fixed point message has been sent to the output channel. This restriction can be relaxed, however, in this case the programmer would have to maintain transitions that check the fixed point condition in every state of each synchroniser along the path.

4.2.0.3 Output from the Serial Replication Network

Now we will clarify how the serial replication network is wired to the rest of the *AstraKahn* application network and how the output is produced. Strictly speaking, the serial replication is not just a wiring pattern since it does not simply wire the replicas of its operand network. It also creates a set of output channels and augments the replicas with some auxiliary synchronisers.

The serial replication N^* defines the output channel set \mathcal{N}_{out} as follows:

$$\mathcal{N}_{out} = \{name(c) \mid c \in \mathcal{O} \wedge fp(c)\}$$

where \mathcal{O} is the output channel set of N and the predicate $fp(c)$ is true on any channel c that has a forward fixed point. The serial replication creates a set of fresh output channels \mathcal{O}^* taking the names from the set \mathcal{N}_{out} . A message that is sent to an inactive replica on any channel c with $name(c) \in \mathcal{N}_{out}$ and which satisfies the fixed point

condition on that channel is immediately transferred to the identically named output channel from \mathcal{O}^* . A network in Fig. 4.3 demonstrates how the output is produced from the serial replication of a network that has a single input and a single output channel.




FIGURE 4.3: Output from the serial replication network

The operand network N has the fixed point condition $P = \bigvee_{j=0}^n p_j$, where p_j is the fixed point condition extracted from the j -th synchroniser ($0 \leq j \leq n$) on the fixed point path of N . In order to check whether a message m on channel x' satisfies the condition p_j , a synchroniser $S^{(j)}$ is inserted before every inactive replica N' . If $p_j(m)$ is true, the synchroniser sends the message m on to the input channel of the next synchroniser $S^{(j+1)}$ to check whether p_{j+1} is satisfied. Otherwise, the message m is sent to the input channel x' of the next replica of N . The listing of the synchroniser $S^{(j)}$, $1 \leq j \leq n-1$ is given in Fig. 4.4. The synchronisers $S^{(0)}$ and $S^{(n)}$ have the same structure, however, $S^{(0)}$ reads messages from the output channel x' of the previous replica and $S^{(n)}$ sends the message that satisfies P to the output of N^* . We shall note that for all j for which $p_j = \bigwedge_{i=0}^k p_j^{(i)}$ the synchroniser $S^{(j)}$ has k transitions that check the conditions $p_j^{(i)}$.

```

synch  $S^{(j)}$  ( $x'^{(j-1)} \mid x'^{(j)}, x'$ )
{
  start {
    on:
       $x'^{(j-1)}.p_j$  {
        send this =>  $x'^{(j)}$ ;
      }
      x.else {
        send this =>  $x'$ ;
      }
  }
}

```

FIGURE 4.4: The synchroniser $S^{(j)}$

The merger M' gathers messages that do not satisfy the fixed point condition from all synchronisers $S^{(j)}$ between two consecutive replicas of N and forwards them to the input channel of the next replica. The merger M gathers messages that satisfy the fixed point condition P and forwards the messages to the output channel x of the serial replication network N^* .

4.2.1 Forward Fixed Point Detection

The existence of a forward fixed point requires the synchronisers that are traversed by the fixed point path to have at least one transition in their start states that accepts the fixed point messages and sends them on unchanged and without producing any other output. Coded in the synchroniser language, the transition that defines the fixed point condition p in channel x is presented in Fig. 4.5.

```
start {
  on:
    x.p {
      send this => out;
    }
  ...
}
```

FIGURE 4.5: The start state of a synchroniser that encodes the fixed point condition p in channel x

Note that there exists no other transition on the channel from the start state with the same structure (a single *send* clause). Otherwise, the synchroniser would have the fixed point condition that is the disjunction of conditions in such transitions. The algorithm in Fig. 4.6 checks if a forward fixed point exists on channel x and extracts the fixed point condition from the synchroniser source code. The algorithm supports the renaming of channel x in the synchroniser. Moreover, it may be the case that the transitions that cause a fixed point in the synchronisers send messages to different output channels. The algorithm detects such a situation; however, the branching of the fixed point path is resolved in the context of the whole network.

The fixed point condition of an *AstraKahn* network is formed by the fixed point conditions of its synchronisers that are traversed by the fixed point path. Networks in *AstraKahn* are represented as graphs, thus the fixed point detection is a graph search problem.

A network graph is a directed multigraph because in *AstraKahn* two nodes are not restricted to be connected with only one edge. The graph has four types of nodes, namely a box, a synchroniser, a merger and a network. The fixed point path may traverse nodes of any type except for boxes. If the path traverses a node that is a network, the network must have a forward fixed point as well.

The fixed point detection algorithm (Fig. 4.7) is based on the depth-first search algorithm with the following considerations:

Require: the abstract syntax tree of the synchroniser program (*synch*), the input label of a channel to test for a forward fixed point (*x*)

Ensure: a dictionary (*a, CondList*), where *a* is the output label of the fixed point channel and *CondList* is the list of atomic fixed point conditions

```

1: function EXTRACT_FP(synch, x)
2:   state  $\leftarrow$  getthestartstatetreefromsynch
3:   CondDict  $\leftarrow$  nil

4:   for each trans in trans_list(state) do
5:     if get_port(trans)  $\neq$  x then
6:       continue  $\triangleright$  the transition reads from another channel
7:     end if
8:     if get_goto(trans)  $\neq$  ('start'  $\vee$   $\emptyset$ ) then
9:       continue  $\triangleright$  the transition does not loop around the state state
10:    end if
11:    if get_assign(trans)  $\neq$   $\emptyset$  then
12:      continue
13:    end if
14:    send  $\leftarrow$  get_send(trans)
15:    if get_msg(send)  $\neq$  this then
16:      continue
17:    end if
18:    cond  $\leftarrow$  get_condition(trans)
19:    if cond is not CondDataMsg  $\vee$  cond is not CondEmpty then
20:      continue  $\triangleright$  the condition cannot be a segmentation mark or an .else
21:    end if
22:    out_port  $\leftarrow$  get_port(send)
23:    if cond  $\in$  CondDict(out_port) then
24:      CondDict(out_port)  $\leftarrow$  CondDict(out_port) .. cond
25:    else
26:      CondDict  $\leftarrow$  CondDict .. (out_port, [cond])
27:    end if
28:  end for

29:  return CondDict
30: end function

```

FIGURE 4.6: Extracting the fixed point condition from a synchroniser (assumes that channel *x* is declared as an input and an output channel of the synchroniser)

- The first and the last nodes of the fixed point path for a particular input channel are known; consequently, only the paths between these two nodes in the graph are traversed
- If the search encounters a box, the traversed path is rejected
- If the search encounters a synchroniser, the fixed point condition of the synchroniser is extracted using the function *extract_fp* in Fig. 4.6. If the synchroniser

has no fixed point condition, the traversed path is rejected. Otherwise, the search continues only for the successors of the node that were detected by *extract_fp*

- If the search encounters a merger, it immediately continues to its only successor
- If the search encounters a node that encapsulates a network, the fixed point detection is run on the network. If no fixed point path exists for the network, the traversed path is rejected.

The fixed point detection algorithm runs only on acyclic networks. The wrap-around wiring makes *AstraKahn* networks cyclic, however, the wrap-around channels cannot carry a fixed point. Therefore, these channels must be filtered before the fixed point detection is run.

4.2.2 Discussion

The approach we have presented relies on the ability of synchronisers to encode some checks of the message content and perform different actions depending on the result of a check. As the analysis in previous sections shows, the construction of an operand network with a complex fixed point condition can be quite complicated. In order to avoid having to construct complicated operand networks, we provide an additional fixed point detection strategy that relies on a special port wiring primitive P that transmits messages immediately from one port to another without storing them. The programmer now has to make sure that the fixed point messages are detected within the operand network⁶ and sent to P . The messages cascade through all the active replicas via a chain of P wires and leave the replication network when they encounter an inactive replica. The example in Fig. 4.8 demonstrates how the approach works.

The operand network A in Fig. 4.8 has a single input port x and two output ports. The output port x is intended for the messages that proceed to the next replica of A in the chain, and the output port x' is an auxiliary port for the messages that are supposed to leave the replication pipeline. The serial replication network A^* has a single input and a single output port both named x . During the compilation, the operand network A is encapsulated into the special network N as shown in Fig. 4.8. The network N inherits all the ports from A and adds the corresponding input port x' . The input and output ports x' of N are connected with the wiring primitive P . The output and the input ports x' of the consequent replicas of N are connected with the wiring primitive P as well. A message that A sends to the output port x' cascades through all the active replicas. Once it has reached an inactive replica, the output channel x of A^* is dynamically wired

⁶It can be done in a box

Require: an operand network graph (*graph*), a channel to test for a forward fixed point (*x*), the first and the last node in the fixed point path (*start*, *end*), the list of fixed point conditions gathered along the path (*cond_list*, optional with the default value *empty_list*)

Ensure: a set of fixed point conditions on the channel

```

1: function DETECT_FFP(graph, x, start, end, cond_list = empty_list)
2:   if start is a box then
3:     return empty_list
4:   end if

5:   if start is a synchroniser then
6:     CondDict  $\leftarrow$  extract_fp(start, x)
7:     if CondDict = nil then
8:       return empty_list
9:     end if
10:    Lists  $\leftarrow$  empty_list
11:    cond_list  $\leftarrow$  cond_list .. start_cond
12:    for each succ_node in succ_nodes(start) do
13:      out_port  $\leftarrow$  get_label(edge(start, succ_node))
14:      if out_port  $\in$  keys(CondDict) then
15:        if start = end then
16:          return cond_list
17:        end if
18:        NewLists  $\leftarrow$  detect_ffp(graph, out_port, succ_node, end, cond_list)
19:        for new_list in NewLists do
20:          Lists  $\leftarrow$  Lists .. new_list       $\triangleright$  the fixed point path branches if
21:        end for
22:      end if
23:    end for
24:    return Lists
25:  end if

26:  if start is a merger then
27:    if start = end then
28:      return cond_list
29:    end if
30:    out_port  $\leftarrow$  get_out_port(start)       $\triangleright$  a merger has a single output port
31:    succ_node  $\leftarrow$  get_succ_node(start)
32:    return detect_ffp(graph, out_port, succ_node, end, cond_list)
33:  end if

34:  if start is a network then
35:    start  $\leftarrow$  get a node that has an input port x
36:    Lists  $\leftarrow$  detect_ffp(graph, x, start, end)
37:    if start = end then
38:      return cond_list .. Lists
39:    end if
40:    return Lists
41:  end if
42: end function

```

FIGURE 4.7: A forward fixed point detection in channel *x* (assumes the network graph is connected)




FIGURE 4.8: The operand network A (top) and a possible implementation of its serial replication A^* (bottom)

to the output port x' of the last active replica of N . When an inactive replica becomes active, the port x' is rewired with the input port of this replica using P .

4.3 Reverse Fixed Point

A reverse fixed point on channel x is a state of a replica, in which it transmits messages from channel x unchanged. A state of a replica is formed by the states of its synchronisers. *AstraKahn* does not analyse boxes and it can determine the operand network behaviour only from its synchronisers. Thus, in order for the operand network to be analysable by *AstraKahn* it must contain a path that does not traverse boxes and which may traverse synchronisers. Every synchroniser that is traversed by the path must have at least one state, in which it accepts a message from channel x unconditionally and sends it on to the next synchroniser along the path without storing or modifying the message.

The existence of a reverse fixed point state requires the operand network to have some topological properties that are formally defined as follows. Consider a network N that has an input and an output channel, both named x .

Definition 2. The network N is said to have a reverse fixed point in x if and only if the following requirements are satisfied:

1. A unique non-branching path from the input to the output channel x exists that does not traverse any boxes
2. Every synchroniser S_i on the path has a subset of states, which we denote as s_i , such that in each of these states every message on the path is immediately

transferred without being changed or stored, causing the synchroniser to remain in the same state⁷. In a state from s_i the synchroniser S_i may still be sensitive to other input channels, as long as this does not, under any circumstances, cause a transition to a state outside s_i

The network N is said to be in a reverse fixed point state on channel x when each S_i is in a state that belongs to its s_i .

4.3.0.1 Rewiring of the Reverse Fixed Point

The reverse fixed point optimises an input connection that has to cascade through the chain to a replica that is ready to accept the data.

The operand network N in Fig. 4.9 has two input and two output channels x and y . Any input channel x wired to an active replica of N that transitions to a reverse fixed point state on that channel is disconnected from the replica and dynamically rewired to the input port x of the next replica in the chain.




FIGURE 4.9: Rewiring of a Reverse Fixed Point replica $N(rfp)$ on channel x

If a replica transitions to a reverse fixed point state on every one of its input channels, no box is running and the channels are empty, the replica is removed from the chain.

4.3.1 Reverse Fixed Point Detection

The existence of a reverse fixed point on channel x requires the synchronisers that are traversed by the fixed point path to have at least one state, in which they unconditionally accept messages from that channel, send them on along the path unchanged and then transition back to the same state. A synchroniser that is in the reverse fixed point state never transitions from it. Coded in the synchroniser language, a transition that makes a state of a synchroniser the reverse fixed point state on channel x is given in Fig. 4.10.

⁷We shall note that the values of any state variables form a part of the synchroniser state

```
s {  
  on:  
    x {  
      send this => out;  
      goto s;  
    }  
    ...  
}
```

FIGURE 4.10: The reverse fixed point state s of a synchroniser on channel x

Note that all other transitions from state s must come back to state s . As long as they do they can even change the state variables of the synchroniser; the reverse fixed point is unconditional and it exists for any values of the store variables.

The algorithm in Fig. 4.11 checks if a reverse fixed point exists on channel x and extracts all fixed point states from the synchroniser source code. The algorithm is designed to work with the detection algorithm in Fig. 4.7.

Require: the abstract syntax tree of the synchroniser program (*synch*), the input label of a channel to test for a forward fixed point (*x*)

Ensure: a dictionary (*a*, *StateList*), where *a* is the output label of the fixed point channel and *StateList* is the list of the reverse fixed point states of the synchroniser

```

1: function EXTRACT_FP(synch, x)
2:   rfp_states  $\leftarrow$  all states from synch that have transitions like in Fig. 4.10
3:   result  $\leftarrow \emptyset$ 

4:   while rfp_states  $\neq$  result do
5:     if result  $\neq \emptyset$  then
6:       rfp_states  $\leftarrow$  result
7:       result  $\leftarrow \emptyset$ 
8:     end if
9:     for each state in rfp_states do
10:      gotos  $\leftarrow$  all destination states in state
11:      if gotos  $\setminus$  rfp_states  $= \emptyset$  then
12:        result  $\leftarrow$  result  $\cup$  state
13:      end if
14:    end for
15:    if result  $= \emptyset$  then
16:      return no reverse fixed point state found
17:    end if
18:  end while

19:  StateDict  $\leftarrow$  nil
20:  for each state in result do
21:    rfp_out  $\leftarrow$  output channel labels of the RPF transitions from state
22:    for each out in rfp_out do
23:      if out  $\notin$  StateDict then
24:        StateDict(out)  $\leftarrow$  state
25:      else
26:        StateDict(out)  $\leftarrow$  StateDict(out) .. state
27:      end if
28:    end for
29:  end for

30:  return StateDict
31: end function

```

FIGURE 4.11: Extracting the reverse fixed point from a synchroniser (assumes that channel *x* is declared as an input and an output channel of the synchroniser)

Chapter 5

Conclusion

5.1 Summary

This thesis presented the implementation of *AstraKahn* synchronisers and their role in the serial replication wiring pattern. It described

Synchronisers are programmed in a dedicated language that is described in [1]. We provide some minor syntax improvements and explain how each language construct should be used in a detailed guide. An *AstraKahn* synchroniser has non-deterministic behaviour. We give an execution algorithm which defines the ordering of non-deterministic choices made by the synchroniser, and is a basis for the synchroniser runtime. We implement the language compiler that generates the data structure to be interpreted by the runtime, and the communication passport of the synchroniser. The compiler performs some semantic and type checking and reports errors in the source code as well.

In *AstraKahn* the output from the serial replication pipeline is defined using the concept of fixed point. A fixed point is a message that is not changed after it has been processed by a newly created replica. In order to detect fixed point messages, *AstraKahn* must be provided with a pattern that matches all of them. We show how this pattern can be embedded into the operand network of the serial replication combinator, so that the programmer does not have to specify it explicitly within the *AstraKahn* application code. Since the original approach to the output from the serial replication is quite complicated, we provide a simple alternative. In order to suppress the growth of the replica chain, *AstraKahn* introduces a reverse fixed point, which is a state of a replica that is known not to change an input connection anymore and thus can be optimised out. We give formal definitions of both kinds of fixed point and provide the algorithms for the *AstraKahn* compiler to detect them.

5.2 Future Work

In this chapter we have presented the language for *AstraKahn* synchronisers. We have developed the language compiler that generates the CAL passport of the synchroniser. The message format in *AstraKahn* is based on the Message Definition Language (MDL) with the restriction that data on streams are organised as collections of alternative records of label-value pairs. A value in a record can have any structure that is allowed by the MDL. The synchroniser that we have presented matches only a top-level structure of a message. The MDL generates a much broader set of terms than the current version synchroniser can synchronise; however, it needs to be elaborated whether the implementation of lower level synchronisation in synchronisers is useful for the real world applications.

The current version of the language does not define flow inheritance in synchronisers. The synchroniser *code* only needs to access the label-value parts of the message it matches. Thus, the flow inheritance should be supported outside of the synchroniser definition and the question how it should be done is left open.

The compiler we have implemented does not optimise the code deeply. We have only implemented an elimination of unused states. However, a broader set of dead code elimination optimisations can be implemented¹.

Serial replication: Implementation of the detection algorithms in the *AstraKahn* compiler.

The long-term goal of the *AstraKahn* project is to provide an environment for development of scalable concurrent applications that do not require manual tuning efforts. Though the adaptive self-regulation mechanism is not known to propose something for synchronisers, we try to review the possible strategies for the pressure propagation through synchroniser.

¹E.g. a conditional transition can be eliminated if its condition always evaluates to *false*

Appendix A

The Syntax of the **AstraKahn** Synchroniser

A.1 Full Grammar

The full grammar of the synchroniser language can be found in Fig. [A.1](#)

The grammar of integer expression used in the synchroniser implementation can be found in Fig. [A.2](#).

A.2 Keywords, Reserved Words and Punctuation

The keywords, the reserved words and the punctuation used in the **AstraKahn** synchroniser syntax are given in Fig. [A.3](#).

A.3 The Abstract Syntax Tree of the Synchroniser

The structure of the abstract syntax tree used in the synchroniser compiler is described in Fig. [A.4](#). Each node is a Node subclass that lists the attributes and child nodes of the class: $\langle name \rangle^*$ - a child node, $\langle name \rangle^{**}$ - a sequence of child nodes and $\langle name \rangle$ - an attribute.

$$\begin{array}{lcl}
\langle \text{int_exp_c} \rangle & ::= & \langle \text{NUMBER} \rangle \\
& | & \langle \text{ID} \rangle \\
& | & \langle ' \rangle \langle \text{int_exp_c} \rangle \langle ' \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '+' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '-' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '*' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '/' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle \% \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '<<' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '>>' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '|' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '\&' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '^' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle '-' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '<' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '>' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '==> \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '!=> \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '<=> \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '>=> \langle \text{int_exp_c} \rangle \\
& | & \langle '!' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '\&\&' \rangle \langle \text{int_exp_c} \rangle \\
& | & \langle \text{int_exp_c} \rangle \langle '||' \rangle \langle \text{int_exp_c} \rangle
\end{array}$$
FIGURE A.2: The syntax of the integer expression in *AstraKahn* synchroniser

Keywords	synch, store, state, int, enum, start, on, elseon, else, do, send, goto
Reserved words	nil, this
Punctuation	braces, brackets, parantheses, the comma, the dot, the semi-colon, the plus sign, the minus sign, the ampersand, the at sign, the question mark, the bar-bar sign, the equality sign, the arrow

FIGURE A.3: *AstraKahn* synchroniser keywords, reserved words and punctuation

```

# inputs -> PortList, outputs -> PortList, decls -> DeclList, states -> StateList
Sync: [name, inputs*, outputs*, decls*, states*]

# ports -> [Port, ...]
PortList: [ports**]

# depth_exp -> ID | NUMBER | DepthExp | DepthNone
Port: [name, depth_exp*]
DepthExp: [depth, shift]
DepthNone: []

# decls -> [StoreVar | StateVar, ...]
DeclList: [decls**]

StoreVar: [name]

# type -> IntType | EnumType
StateVar: [name, type*]
IntType: [size]

# labels -> [ID, ...]
EnumType: [labels**]

# states -> [State, ...]
StateList: [states**]

# trans_orders -> [TransOrder, ...]
State: [name, trans_orders**]

# trans_stmt -> [Trans, ...]
TransOrder: [trans_stmt**]

# condition -> CondSegmark | CondDataMsg | CondEmpty | CondElse
# guard -> IntExp
# actions -> [Assign | Send | Goto, ...]
Trans: [port, condition*, guard*, actions**]

CondSegmark: [depth]

# labels -> [ID, ...]
CondDataMsg: [choice, labels**, tail]

CondEmpty: []
CondElse: []

# rhs -> DataExp | IntExp
Assign: [lhs, rhs*]

# items -> [ItemThis | ItemVar | ItemExpand | ItemPair, ...]
DataExp: [items**]
ItemThis: []
ItemVar: [name]
ItemExpand: [name]
# value -> ID | IntExp
ItemPair: [label, value*]

# msg -> MsgSegmark | MsgData | MsgNil
Send: [msg*, port]

# depth -> ID | IntExp
MsgSegmark: [depth*]

# data_exp -> DataExp
MsgData: [choice, data_exp*]
MsgNil: []

# states -> [ID, ...]
Goto: [states**]

ID: [name]
NUMBER: [value]
IntExp: [exp]

```

FIGURE A.4: The abstract syntax tree of a synchroniser

Bibliography

- [1] Alex Shafarenko. Astrakahn: A coordination language for streaming networks. *CoRR*, abs/1306.6029, 2013.
- [2] George A. Papadopoulos and Farhad Arbab. Coordination models and languages. In *ADVANCES IN COMPUTERS*, pages 329–400. Academic Press, 1998.
- [3] Sudhir Ahuja, Nicholas Carriero, and David Gelernter. Linda and friends. *Computer*, 19(8):26–34, August 1986.
- [4] Message Passing Interface Forum. MPI: A Message-Passing Interface Standard Version 3.0, 2012. Chapter author for Collective Communication, Process Topologies, and One Sided Communications.
- [5] Farhad Arbab. Reo: a channel-based coordination model for component composition. *Mathematical Structures in Computer Science*, 14(3):329–366, 2004.
- [6] Christel Baier, Marjan Sirjani, Farhad Arbab, and Jan Rutten. Modeling component connectors in reo by constraint automata. *Sci. Comput. Program.*, 61(2):75–113, July 2006.
- [7] Robert Stephens. A survey of stream processing. *Acta Inf.*, 34(7):491–541, 1997.
- [8] Edward A. Ashcroft and William W. Wadge. Lucid, a nonprocedural language with iteration. *Commun. ACM*, 20(7):519–526, 1977.
- [9] Gilles Kahn. The semantics of simple language for parallel programming. In *IFIP Congress*, pages 471–475, 1974.
- [10] Edward Ashford Lee and David G. Messerschmitt. Static scheduling of synchronous data flow programs for digital signal processing. *IEEE Trans. Comput.*, 36(1):24–35, January 1987.

- [11] William Thies, Michal Karczmarek, and Saman P. Amarasinghe. Streamit: A language for streaming applications. In *Proceedings of the 11th International Conference on Compiler Construction*, CC '02, pages 179–196, London, UK, UK, 2002. Springer-Verlag.
- [12] Clemens Grelck, Sven-Bodo Scholz, and Alexander V. Shafarenko. A Gentle Introduction to S-Net: Typed Stream Processing and Declarative Coordination of Asynchronous Components. *Parallel Processing Letters*, 18(2):221–237, 2008.
- [13] Vu Thien Nga Nguyen and Raimund Kirner. Demand-based scheduling priorities for performance optimisation of stream programs on parallel platforms. In *ICA3PP (1)*, pages 357–369, 2013.
- [14] David Beazly et al. Ply (python lex-yacc). <http://www.dabeaz.com/ply/>, February 2011. [Online; accessed 9-December-2014].
- [15] Alfred V. Aho, Monica S. Lam, Ravi Sethi, and Jeffrey D. Ullman. *Compilers: Principles, Techniques, and Tools (2nd Edition)*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2006.